

binaryClassifiers

April 10, 2019

1 Loan Application Approval Classifier

1.0.1 The Data

- `ds-credit.tsv` - File containing information about a customer's credit history.
- `ds-app.tsv` - File containing information regarding the customer's loan application.
- `ds-borrower.tsv` - File containing information regarding the customer's financial details.
- `ds-result.tsv` - File containing the **actual** results of the loan application process.

1.0.2 The Schema

The four files named above are our data sources. We can assume them to be **database tables** with schemas as illustrated below.

Clearly, the feature `CustomerId` is common across all these tables and is the **key** ID that ties them together. Since I am using Python for this project, I am going to combine these tables into a single table, before dividing my data into **training** and **testing** sets for my classifier.

Before doing that however, we must read the files into Python. These files are **tab separated** files, opening them up in a text editor lets me know that they do not have headers. I assigned the headers from the project description to the relevant files in the text editor itself.

Loading the libraries

```
In [74]: import os
import csv
import time
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from prettytable import PrettyTable
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
plt.rcParams['figure.figsize'] = [16, 10]
plt.figure(figsize=(16,10))
# plt.figure(figsize=(16, 10))
from IPython.core.display import display,HTML
display(HTML('<style>.prompt{width: 0px; min-width: 0px; visibility: collapse}</style>'))
```

<IPython.core.display.HTML object>

<Figure size 1152x720 with 0 Axes>

Setting the path

```
In [75]: path = os.getcwd()
         applicationFile = path + '/ds-app.tsv'
         borrowerFile = path + '/ds-borrower.tsv'
         creditFile = path + '/ds-credit.tsv'
         resultFile = path + '/ds-result.tsv'
```

Data Manipulation Let us use the pandas library to read the application data into a dataframe

```
In [76]: applicationdf = pd.read_csv(applicationFile, sep='\t', header = 0)
         applicationdf.head()
```

```
Out [76]:
```

					CustomerID	LoanPayoffPeriodInMonths	LoanReason	\
0	741	36	goods	9566	2	none	NaN	NaN
1	439	12	goods	2279	4	none	NaN	NaN
2	501	30	auto	3857	4	none	NaN	NaN
3		147	12	auto	900	4	none	NaN
4	821	6	goods	1898	1	none	NaN	NaN

	RequestedAmount	InterestRate	Co-Applicant
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

There are 2 interesting things to note - 1. The data is **not** tab separated, but **space** separated as all the data have been incorrectly parsed into a single column. 2. The headers which I manually input into the files, correctly parse into separate columns.

To handle this problem - * I will use the csv library to deal with the data row by row. * This will lead to a row-based data structure, which is not conducive for data analysis. I will apply additional transformations to change the data to a column-based dataframe. * It is safe to assume that the other datasets also face the same problem. To avoid checking those files individually and to save time, I will apply my solution uniformly across each file.

```
In [77]: files = [applicationFile, borrowerFile, creditFile, resultFile]
         holder = {}
         for file in files:
             with open(file) as fp:
                 reader = csv.reader(fp, delimiter='\t')
                 header = next(reader, None)
```

```

table = []
for row in reader:
    elements = row[0].split()
    table.append(elements)
df = pd.DataFrame(table, columns = header)
holder[file] = df
del elements
del df
del table

```

In the code above, I transform each data table from a row-oriented structure to a column based pandas dataframe. I store those pandas dataframes in a temporary dictionary, indexed by the file names. I can now assign each pandas dataframe to its own variable and take a deeper dive into each data table.

```

In [78]: applicationdf = holder[applicationFile]
        borrowerdf = holder[borrowerFile]
        creditdf = holder[creditFile]
        resultdf = holder[resultFile]
        del holder

```

We can now check if the correct data has been assigned to the correct variables.

```

In [79]: applicationdf.head()

```

```

Out [79]:   CustomerID  LoanPayoffPeriodInMonths  LoanReason  RequestedAmount  InterestRate  \
0         741                36      goods           9566             2
1         439                12      goods           2279             4
2         501                30      auto           3857             4
3         147                12      auto            900             4
4         821                 6      goods           1898             1

      Co-Applicant
0         none
1         none
2         none
3         none
4         none

```

```

In [80]: borrowerdf.head()

```

```

Out [80]:   CustomerID  YearsAtCurrentEmployer  YearsInCurrentResidence  Age  RentOrOwnHome  \
0         765                1                1  46             rent
1         668                4                4  49             owned
2          68               10+                4  57             rent
3         805                4                2  41             owned
4         495                7                4  24             rent

      TypeOfCurrentEmployment  NumberOfDependantsIncludingSelf

```

0	skilled	1
1	skilled	1
2	skilled	1
3	skilled	1
4	skilled	1

```
In [81]: creditdf.head()
```

```
Out[81]:
```

	CustomerID	CheckingAccountBalance	DebtsPaid	SavingsAccountBalance	\
0	374	debt	paid	some	
1	346	some	delayed	none	
2	345	some	paid	none	
3	243	some	delayed	some	
4	662	debt	delayed	some	

	CurrentOpenLoanApplications
0	2
1	1
2	1
3	1
4	2

```
In [82]: resultdf.head()
```

```
Out[82]:
```

	CustomerID	WasTheLoanApproved
0	1	Y
1	2	Y
2	3	N
3	4	Y
4	5	Y

Now, we must combine or **join** these dataframes into a single dataframe, before we do that however, we must ensure that each of them has the same length. Doing so would bolster the following assumptions - * Each customer has only 1 row in each data table. * The customers are uniform across each data table.

If the assumptions hold, we can simply combine the constituent tables into 1 big dataset.

```
In [83]: print ("application   : ", applicationdf.shape)
          print ("borrower     : ", borrowerdf.shape)
          print ("credit       : ", creditdf.shape)
          print ("result      : ", resultdf.shape)
```

```
application   : (748, 6)
borrower      : (749, 7)
credit        : (749, 5)
result        : (750, 2)
```

The results above tell us - * There are more results than applications => There are redundant customers in the results. * There are more results than borrower and credit. This bolsters our above claim and we will have to remove these redundant customers from our analysis. * There is at least one customer, for whom we do not have the application details but do have borrower and credit information. Now, without knowing important application features like the loan amount, duration or interest rate, we cannot make an accurate prediction as to whether they should be approved for a loan. We will have to find and remove this customer from our analysis. * Conversely, it could be the case that 1 customer has 2 separate rows for borrower and application. In that case, we will have to decide which of the rows could play a greater role in our analysis and act accordingly.

```
In [84]: applicationdf.set_index('CustomerID', inplace=True)
         borrowerdf.set_index('CustomerID', inplace=True)
         creditdf.set_index('CustomerID', inplace=True)
         resultdf.set_index('CustomerID', inplace=True)
```

We set CustomerID as the index across all the dataframes and then use the `index.difference` method to check which CustomerID are absent from the results but present in the other dataframes. These CustomerID values will form our **blind test data**.

```
In [85]: applicationdf.index.difference(resultdf.index)
```

```
Out [85]: Index(['751', '752', '753', '754', '755', '757', '758', '759', '760', '761',
               '762', '763', '765', '766', '767', '768', '769', '770', '771', '772',
               '773', '774', '775', '776', '777', '778', '780', '781', '782', '784',
               '785', '786', '788', '790', '791', '794', '795', '796', '797', '798',
               '799', '800', '801', '802', '803', '804', '805', '807', '810', '811',
               '812', '813', '814', '815', '816', '817', '818', '819', '820', '821',
               '822', '823', '824', '825', '827', '828', '829', '830', '831', '832',
               '833', '835', '836', '837', '838', '839', '840', '841', '842', '843',
               '844', '845', '846', '847', '848'],
              dtype='object', name='CustomerID')
```

```
In [86]: borrowerdf.index.difference(resultdf.index)
```

```
Out [86]: Index(['751', '753', '754', '755', '756', '757', '758', '759', '760', '761',
               '762', '763', '764', '765', '766', '767', '768', '769', '770', '771',
               '772', '773', '774', '775', '776', '777', '778', '779', '780', '781',
               '782', '783', '784', '785', '786', '787', '788', '789', '790', '792',
               '793', '795', '796', '797', '799', '800', '801', '802', '803', '804',
               '805', '806', '807', '808', '809', '810', '811', '812', '813', '814',
               '815', '816', '817', '818', '819', '821', '822', '823', '824', '825',
               '826', '827', '828', '829', '830', '831', '832', '833', '835', '836',
               '837', '839', '840', '841', '842', '845', '846', '847', '848', '849'],
              dtype='object', name='CustomerID')
```

```
In [87]: creditdf.index.difference(resultdf.index)
```

```
Out[87]: Index(['751', '753', '754', '755', '756', '758', '759', '761', '762', '764',
              '765', '766', '767', '768', '769', '770', '771', '772', '773', '775',
              '776', '777', '778', '779', '780', '781', '782', '783', '784', '785',
              '786', '787', '788', '789', '790', '791', '792', '793', '794', '795',
              '796', '797', '798', '799', '800', '801', '802', '804', '806', '807',
              '808', '809', '810', '811', '812', '813', '814', '816', '817', '818',
              '819', '821', '822', '823', '824', '825', '826', '827', '829', '830',
              '831', '832', '833', '834', '835', '836', '837', '838', '839', '840',
              '841', '842', '843', '844', '845', '846', '847', '848', '849'],
              dtype='object', name='CustomerID')
```

As we can see, the customerID values from 751 onwards do not have results. These form our blind test data. Let us now join these data frames together. We perform these steps to join the data - * **Left outer join** application data with borrower data, this ensures that we retain only customers who have valid applications. * **Left outer join** the above table with credit data. * **Left outer join** the above table with the result data.

```
In [88]: join1 = pd.merge(applicationdf,
                          borrowerdf,
                          right_index=True,
                          left_index=True,
                          how = 'left')

join1.shape
```

```
Out[88]: (748, 11)
```

```
In [89]: join2 = pd.merge(join1,
                          creditdf,
                          right_index=True,
                          left_index=True,
                          how = 'left')

join2.shape
```

```
Out[89]: (748, 15)
```

```
In [90]: join3 = pd.merge(join2,
                          resultdf,
                          right_index=True,
                          left_index=True,
                          how = 'left')

join3.shape
```

```
Out[90]: (748, 16)
```

```
In [91]: join3.index = join3.index.map(int)
join3.sort_index(inplace=True)
join3.head()
```

Out [91]:

CustomerID	LoanPayoffPeriodInMonths	LoanReason	RequestedAmount	InterestRate	\
2	9	goods	3074	1	
3	12	auto	939	4	
4	9	auto	2507	2	
5	18	goods	2238	2	
6	24	repairs	5507	3	

CustomerID	Co-Applicant	YearsAtCurrentEmployer	YearsInCurrentResidence	Age	\
2	none	4	2	33	
3	none	7	2	28	
4	none	10+	4	51	
5	none	4	1	25	
6	none	10+	4	44	

CustomerID	RentOrOwnHome	TypeOfCurrentEmployment	\
2	owned	skilled	
3	owned	skilled	
4	free	unskill	
5	owned	skilled	
6	free	skilled	

CustomerID	NumberOfDependantsIncludingSelf	CheckingAccountBalance	DebtsPaid	\
2	2	none	delayed	
3	1	high	delayed	
4	1	none	paid	
5	1	none	delayed	
6	1	none	delayed	

CustomerID	SavingsAccountBalance	CurrentOpenLoanApplications	\
2	none	2	
3	high	3	
4	high	1	
5	some	2	
6	some	2	

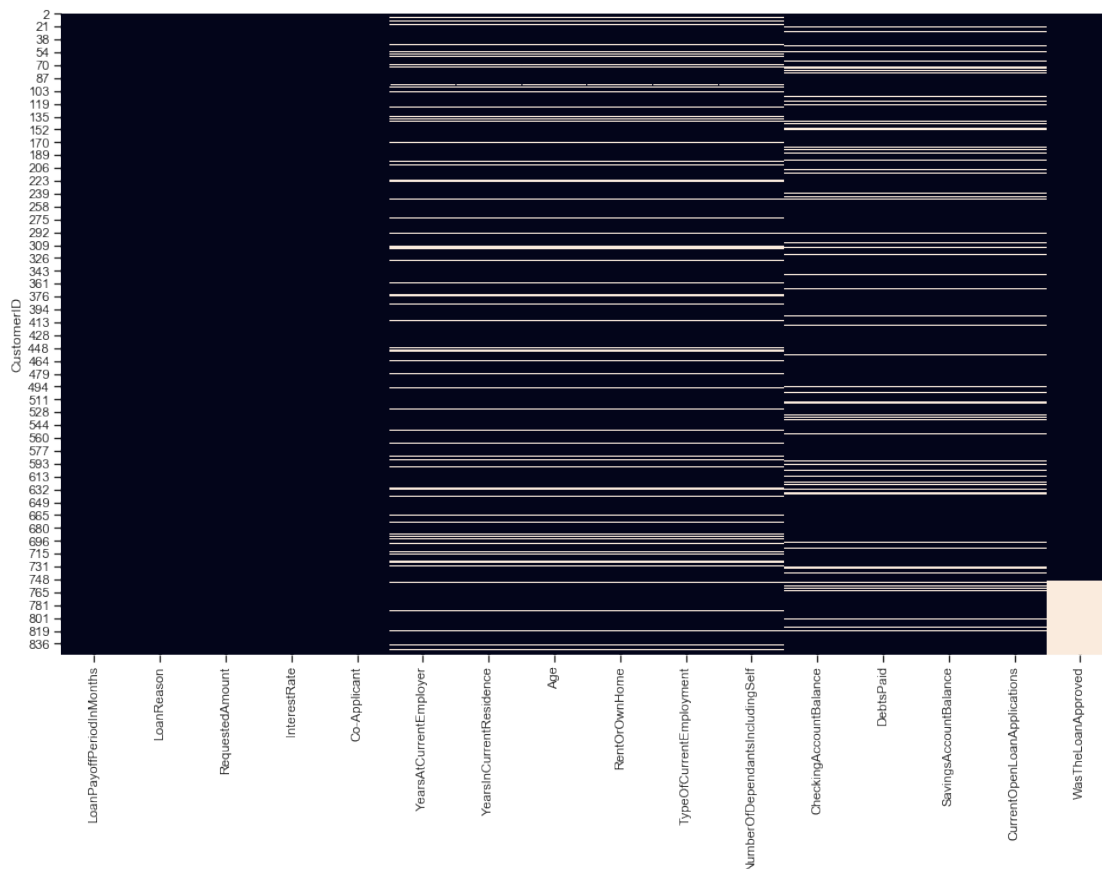
CustomerID	WasTheLoanApproved
2	Y
3	N
4	Y
5	Y
6	Y

1.1 Exploratory Data Analysis

Now that we have our data in one combined dataset, we still have to - * Split our data into **training**, **testing** and **blind test** datasets. * Explore our data visually to find - * Missing Values. * Data imbalance. * Visual patterns that show relationships.

```
In [92]: sns.heatmap(join3.isnull(), cbar=False)
```

```
Out[92]: <matplotlib.axes._subplots.AxesSubplot at 0x168a81b30b8>
```



The plot above shows us missing values in the joined dataset. Let us look at WasTheLoanApproved, it has a solid chunk of missing values towards the end of the dataset. This chunk forms our **blind test data**.

Before we split our data, let us ensure that the data types are as expected

```
In [93]: join3.dtypes
```

```
Out[93]: LoanPayoffPeriodInMonths    object
         LoanReason                  object
         RequestedAmount              object
         InterestRate                object
         Co-Applicant                 object
```


YearsAtCurrentEmployer	object
YearsInCurrentResidence	object
Age	object
RentOrOwnHome	object
TypeOfCurrentEmployment	object
NumberOfDependantsIncludingSelf	object
CheckingAccountBalance	object
DebtsPaid	object
SavingsAccountBalance	object
CurrentOpenLoanApplications	object
WasTheLoanApproved	object
dtype:	object

We need to map these values to the correct datatype. The columns that are numeric need to be explicitly converted to numeric, we can leave the columns that are strings alone.

```
In [94]: join3[["LoanPayoffPeriodInMonths", "RequestedAmount", "InterestRate",
               "YearsInCurrentResidence", "Age", "NumberOfDependantsIncludingSelf",
               "CurrentOpenLoanApplications"]] = join3[["LoanPayoffPeriodInMonths", "RequestedAmount",
               "YearsInCurrentResidence", "Age", "NumberOfDependantsIncludingSelf",
               "CurrentOpenLoanApplications"]].apply(
               lambda x: pd.to_numeric(x, errors='coerce'), axis=0)
join3.dtypes
```

```
Out [94]: LoanPayoffPeriodInMonths    int64
LoanReason                          object
RequestedAmount                     int64
InterestRate                        int64
Co-Applicant                        object
YearsAtCurrentEmployer              object
YearsInCurrentResidence              float64
Age                                 float64
RentOrOwnHome                       object
TypeOfCurrentEmployment              object
NumberOfDependantsIncludingSelf        float64
CheckingAccountBalance               object
DebtsPaid                           object
SavingsAccountBalance               object
CurrentOpenLoanApplications          float64
WasTheLoanApproved                   object
dtype: object
```

```
In [95]: blindTestdf = join3[join3['WasTheLoanApproved'].isnull()]
blindTestdf.shape
```

```
Out [95]: (85, 16)
```

```
In [96]: combineddf = join3[-join3['WasTheLoanApproved'].isnull()]
combineddf.shape
```

```
Out [96]: (663, 16)
```

Our blind test data has 85 datapoints, now, ideally we would like a 75 : 25 training-test split. Keep in mind that we now have a total of 663 data points. We will have to split this data to get the **training** and **testing** datasets. We will train our model on the training data, test it on our testing dataset and make predictions for the blind dataset.

```
In [97]: del join1
         del join2
         del join3
```

```
In [98]: from sklearn.model_selection import train_test_split
```

```
# y = combineddf['WasTheLoanApproved']
# x_train, x_test, y_train, y_test = train_test_split(combineddf, y, test_size = 0.25)
# print(x_train.shape, y_train.shape)
# print(x_test.shape, y_test.shape)
```

```
train, test = train_test_split(combineddf, test_size = 0.25)
print(train.shape, test.shape)
```

```
(497, 16) (166, 16)
```

1.1.1 Handling missing data

Let us now explore the missing data values in our dataset. First let us look at what fraction of our data is missing.

```
In [99]: round(train.isnull().sum()/train.shape[0] * 100, 2)
```

```
Out[99]: LoanPayoffPeriodInMonths      0.00
         LoanReason                     0.00
         RequestedAmount                0.00
         InterestRate                  0.00
         Co-Applicant                   0.00
         YearsAtCurrentEmployer        11.47
         YearsInCurrentResidence       11.47
         Age                           11.47
         RentOrOwnHome                 11.47
         TypeOfCurrentEmployment       11.47
         NumberOfDependantsIncludingSelf  11.47
         CheckingAccountBalance        10.66
         DebtsPaid                     10.66
         SavingsAccountBalance         10.66
         CurrentOpenLoanApplications  10.66
         WasTheLoanApproved             0.00
         dtype: float64
```

```
In [100]: round(test.isnull().sum()/test.shape[0] * 100, 2)
```

```
Out[100]: LoanPayoffPeriodInMonths      0.00
          LoanReason                     0.00
          RequestedAmount                 0.00
          InterestRate                    0.00
          Co-Applicant                    0.00
          YearsAtCurrentEmployer          12.65
          YearsInCurrentResidence          12.65
          Age                             12.65
          RentOrOwnHome                    12.65
          TypeOfCurrentEmployment          12.65
          NumberOfDependantsIncludingSelf    12.65
          CheckingAccountBalance           15.06
          DebtsPaid                        15.06
          SavingsAccountBalance            15.06
          CurrentOpenLoanApplications      15.06
          WasTheLoanApproved               0.00
          dtype: float64
```

```
In [101]: round(blindTestdf.isnull().sum()/blindTestdf.shape[0] * 100, 2)
```

```
Out[101]: LoanPayoffPeriodInMonths      0.00
          LoanReason                     0.00
          RequestedAmount                 0.00
          InterestRate                    0.00
          Co-Applicant                    0.00
          YearsAtCurrentEmployer          9.41
          YearsInCurrentResidence          9.41
          Age                             9.41
          RentOrOwnHome                    9.41
          TypeOfCurrentEmployment          9.41
          NumberOfDependantsIncludingSelf    9.41
          CheckingAccountBalance           11.76
          DebtsPaid                        11.76
          SavingsAccountBalance            11.76
          CurrentOpenLoanApplications      11.76
          WasTheLoanApproved              100.00
          dtype: float64
```

In both our data sets, the %age of missing values ranges from 9% – 12%. We can handle missing values in the following ways - * **The easy way** - Remove them from our data. * **The proper way** - Impute the mean values for numerical data and most common value for string data.

At 9% – 12%, and with the few datapoints that we have, we cannot afford to remove the missing data if we wish to make a strong classifier. We separate the data by numeric and object type and then apply our strategies. Note, the datasets were joined in a way that only valid applications would be considered, therefore, the columns in applicationdf do not have missing values, we can exclude these columns from the missing data imputation process.

```
In [102]: num_dtypes = [i for i in train.select_dtypes(include = np.number).columns if i not in
                      obj_dtypes = [i for i in train.select_dtypes(include = np.object).columns if i not in
                      obj_dtypes = [x for x in obj_dtypes if x != 'WasTheLoanApproved']
```

```
In [103]: num_dtypes
```

```
Out[103]: ['YearsInCurrentResidence',  
           'Age',  
           'NumberOfDependantsIncludingSelf',  
           'CurrentOpenLoanApplications']
```

```
In [104]: obj_dtypes
```

```
Out[104]: ['YearsAtCurrentEmployer',  
           'RentOrOwnHome',  
           'TypeOfCurrentEmployment',  
           'CheckingAccountBalance',  
           'DebtsPaid',  
           'SavingsAccountBalance']
```

We handle missing data as follows - * For numerical features we impute the **median** value of the data to the missing values. * For strings, we impute the most common value (**mode**) of the data to the missing values.

Let us handle the numerical data types first.

```
In [105]: table = PrettyTable()  
          table.field_names = ["Feature", "Mean", "Median"]  
          for col in num_dtypes:  
              table.add_row([col, round(train[col].mean(), 3), train[col].median()])  
          print(table)
```

Feature	Mean	Median
YearsInCurrentResidence	2.795	3.0
Age	35.677	34.0
NumberOfDependantsIncludingSelf	1.155	1.0
CurrentOpenLoanApplications	1.428	1.0

Since the median value is quite close to the mean, and conveniently, a whole number, we will use the it to impute the missing values for numerical features. We will also create 2 new dataframes to hold our train and test sets.

```
In [106]: train_new = train[[i for i in train.columns if i not in num_dtypes + obj_dtypes]]  
          test_new = test[[i for i in test.columns if i not in num_dtypes + obj_dtypes]]  
          blind_new = blindTestdf[[i for i in blindTestdf.columns if i not in num_dtypes + obj_dtypes]]
```

```
In [107]: for j, col in enumerate(num_dtypes):  
            train_new[j] = train[col].fillna(int(train[col].median()))  
            train_new = train_new.rename(columns = {j: col})  
            test_new[j] = test[col].fillna(int(test[col].median()))  
            test_new = test_new.rename(columns = {j: col})  
            blind_new[j] = blindTestdf[col].fillna(int(blindTestdf[col].median()))  
            blind_new = blind_new.rename(columns = {j: col})
```

We will now handle the features with string/object data type.

```
In [108]: for k, col in enumerate(obj_dtypes):
          train_new[k] = train[col].fillna(train[col].mode()[0])
          train_new = train_new.rename(columns={k: col})
          test_new[k] = test[col].fillna(test[col].mode()[0])
          test_new = test_new.rename(columns={k: col})
          blind_new[j] = blindTestdf[col].fillna(blindTestdf[col].mode()[0])
          blind_new = blind_new.rename(columns = {j: col})
```

```
In [109]: train_new.isnull().sum()
```

```
Out[109]: LoanPayoffPeriodInMonths      0
          LoanReason                     0
          RequestedAmount                 0
          InterestRate                   0
          Co-Applicant                   0
          WasTheLoanApproved              0
          YearsInCurrentResidence         0
          Age                             0
          NumberOfDependantsIncludingSelf   0
          CurrentOpenLoanApplications     0
          YearsAtCurrentEmployer          0
          RentOrOwnHome                   0
          TypeOfCurrentEmployment         0
          CheckingAccountBalance          0
          DebtsPaid                       0
          SavingsAccountBalance           0
          dtype: int64
```

```
In [110]: test_new.isnull().sum()
```

```
Out[110]: LoanPayoffPeriodInMonths      0
          LoanReason                     0
          RequestedAmount                 0
          InterestRate                   0
          Co-Applicant                   0
          WasTheLoanApproved              0
          YearsInCurrentResidence         0
          Age                             0
          NumberOfDependantsIncludingSelf   0
          CurrentOpenLoanApplications     0
          YearsAtCurrentEmployer          0
          RentOrOwnHome                   0
          TypeOfCurrentEmployment         0
          CheckingAccountBalance          0
          DebtsPaid                       0
          SavingsAccountBalance           0
          dtype: int64
```

```
In [111]: blind_new.isnull().sum()
```

```
Out [111]: LoanPayoffPeriodInMonths      0
           LoanReason                     0
           RequestedAmount                0
           InterestRate                   0
           Co-Applicant                   0
           WasTheLoanApproved             85
           YearsInCurrentResidence        0
           Age                           0
           NumberOfDependantsIncludingSelf  0
           CurrentOpenLoanApplications    0
           YearsAtCurrentEmployer         0
           RentOrOwnHome                   0
           TypeOfCurrentEmployment        0
           CheckingAccountBalance         0
           DebtsPaid                       0
           SavingsAccountBalance          0
           dtype: int64
```

For the blind test data, we replace the Nans in WasTheLoanApproved with 0. That will help us pass the data to the classifier model, which will then make the predictions for what it believes the values should be for WasTheLoanApproved.

```
In [112]: blind_new["WasTheLoanApproved"].fillna(0, inplace=True)
           blind_new.head()
```

```
Out [112]:      LoanPayoffPeriodInMonths  LoanReason  RequestedAmount  \
CustomerID
751                12      goods           804
752                11      auto          3939
753                24      goods          2828
754                48     busin          3844
755                 9      goods          1126
```

```
      InterestRate  Co-Applicant  WasTheLoanApproved  \
CustomerID
751                4         none                0
752                1         none                0
753                4         none                0
754                4         none                0
755                2         none                0
```

```
      YearsInCurrentResidence  Age  NumberOfDependantsIncludingSelf  \
CustomerID
751                4.0  38.0                1.0
752                4.0  35.0                1.0
753                4.0  22.0                1.0
754                4.0  34.0                2.0
```

755	4.0	49.0	1.0
-----	-----	------	-----

	CurrentOpenLoanApplications	YearsAtCurrentEmployer	RentOrOwnHome	\
CustomerID				
751	1.0	10+	owned	
752	1.0	4	owned	
753	1.0	4	owned	
754	1.0	7	free	
755	1.0	10+	owned	

	TypeOfCurrentEmployment	CheckingAccountBalance	DebtsPaid	\
CustomerID				
751	skilled	none	paid	
752	skilled	none	paid	
753	skilled	debt	paid	
754	unskill	some	paid	
755	skilled	high	paid	

	SavingsAccountBalance
CustomerID	
751	some
752	some
753	high
754	medium
755	medium

We have successfully *approximated* the missing values. Let us now check if the data is **unbalanced**. Class imbalance influences a learning algorithm during training by making the decision rule biased towards the majority class by implicitly learning a model that optimizes the predictions based on the majority class in the dataset. As a result, we'll explore methods to handle the class imbalance problem.

```
In [113]: combineddf.drop(combineddf[combineddf.WasTheLoanApproved == 1].index, inplace=True)
approved = combineddf[combineddf["WasTheLoanApproved"] == 'Y'].shape[0]
denied = combineddf[combineddf["WasTheLoanApproved"] == 'N'].shape[0]
print("Positive examples : ", approved)
print("Negative examples : ", denied)
print("Proportion of negative to positive examples : ", round(denied/approved, 2) * 100)
print("Proportion of negative to total examples : ", round(denied/combineddf.shape[0], 2) * 100)
sns.countplot(combineddf["WasTheLoanApproved"])
plt.xticks((0, 1), ["Loan Approved", "Loan Denied"])
plt.xlabel("")
plt.ylabel("Count")
plt.title("Class counts", y=1, fontdict={"fontsize": 20})
```

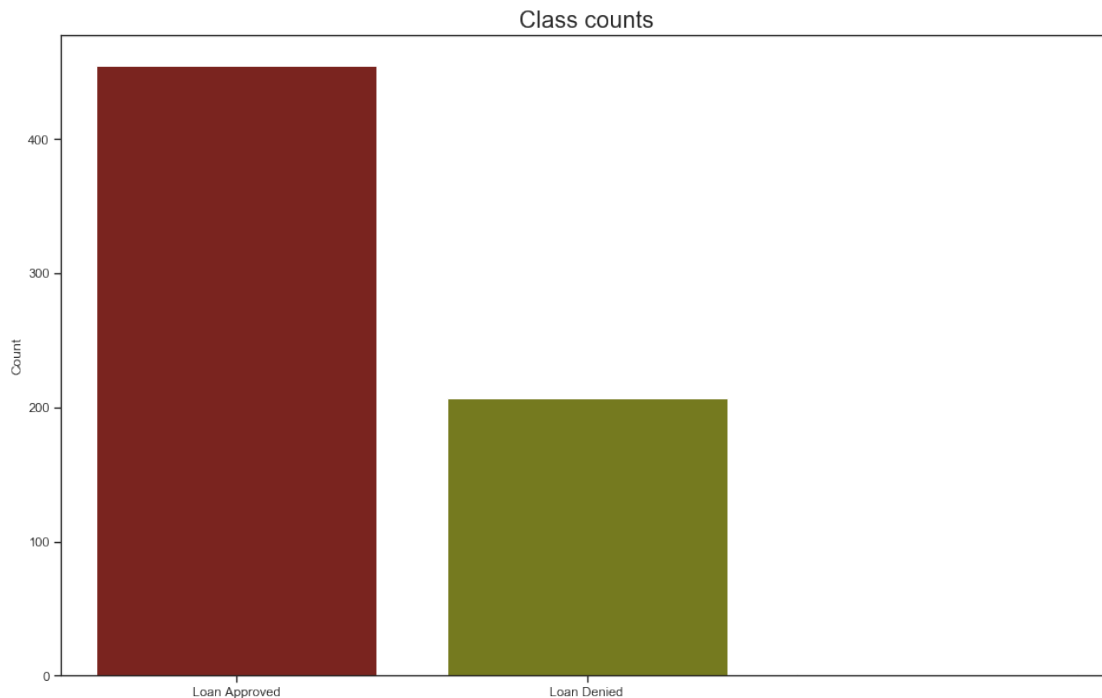
Positive examples : 455

Negative examples : 207

Proportion of negative to positive examples : 45.0 %

Proportion of negative to total examples : 31.0 %

```
Out[113]: Text(0.5, 1, 'Class counts')
```



Classification problems in most real world applications have imbalanced data sets. As we can see, there exists an imbalance between the positive (loan approved) examples (majority class) and the negative (loan denied) examples (minority class). In our example, the negative examples form only 31% of the total examples. Therefore, accuracy is no longer a good measure of performance for different models because if we simply predict all examples to belong to the positive class, we achieve 69% accuracy. Therefore, better metrics for imbalanced data sets are **AUC** (area under the ROC curve) and **f1-score**. Since, the dataset is imbalanced, but not too skewed, we will make our classifiers without using sampling methods.

1.1.2 Visualizing our Data

Now that we have handled missing data, we can now get started with the core of our task, building and evaluating the model. To assist in that process, let us explore our data visually to get a deeper understanding of our features and how they relate to our target variable `WasTheLoanApproved`.

```
In [114]: sns.set(rc={'figure.figsize':(16, 9)})
          sns.set_style("ticks")
          # sns.set_palette("dark")
          # sns.set_palette("Paired")
          sns.set_palette(sns.hls_palette(6, l=.3, s=.8))
```

The above snippet of code sets our parameters for seaborn plots.


```
In [115]: train_new.WasTheLoanApproved.replace(('Y', 'N'), (1, 0), inplace=True)
          test_new.WasTheLoanApproved.replace(('Y', 'N'), (1, 0), inplace=True)

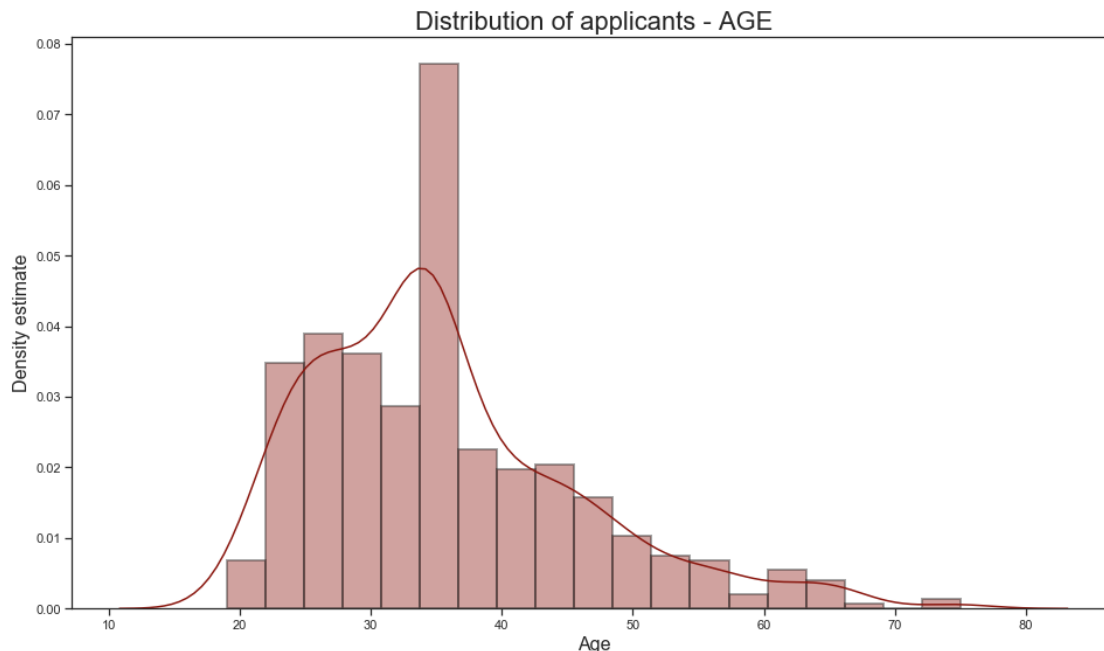
          train_new[['WasTheLoanApproved']] = train_new[['WasTheLoanApproved']].apply(pd.to_numeric)
          test_new[['WasTheLoanApproved']] = test_new[['WasTheLoanApproved']].apply(pd.to_numeric)

          # blind_new.WasTheLoanApproved.replace(('Y', 'N'), (1, 0), inplace=True)
          # blind_new[['WasTheLoanApproved']] = blind_new[['WasTheLoanApproved']].apply(pd.to_numeric)
```

Here are some questions that come to mind when we look at the loan data - * **Who?** - What generally describes the people applying for a loan? * **Why?** - What is driving them to take apply for a loan? * **Amount?** - In general, how much are the applicants looking to borrow and at what interest rates? * **Time?** - What is, in general, the duration of the loan payment period? * **Should we?** - Do they pose a risk of failing to pay their loan? * And finally, are there patterns that indicate which factors contribute to the decision of whether to approve or deny an applicant?

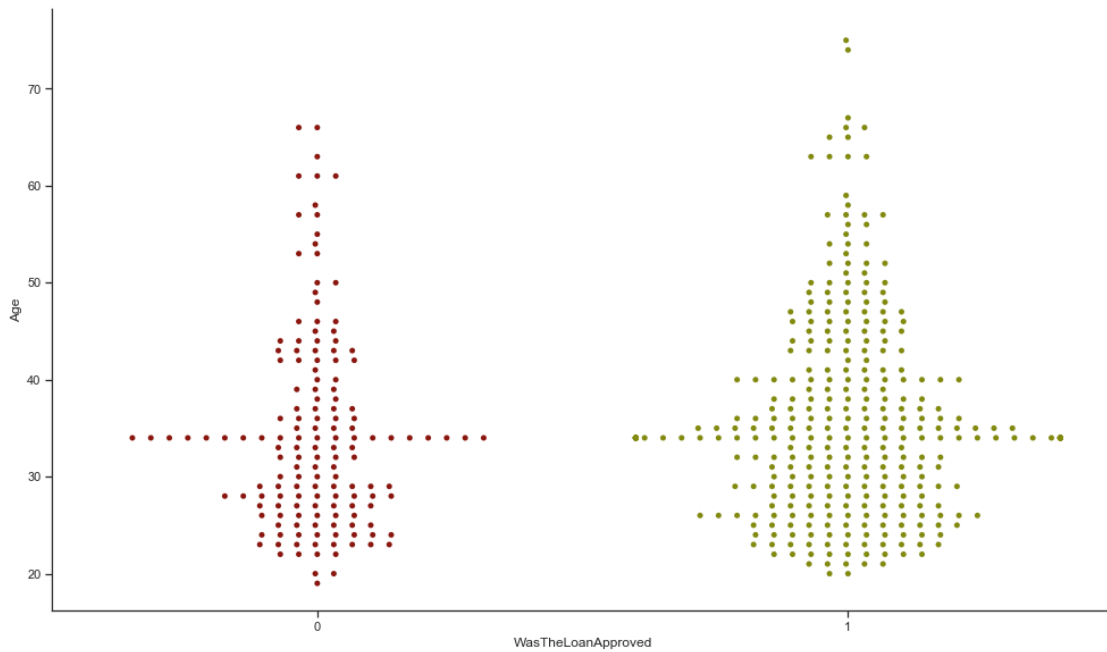
All these questions are very general, but we can get a high level answer to these questions by making some plots. For example, to get an idea of who is applying for these loans, we can explore * The distribution of their ages. * Whether they are applying alone. * Whether they rent/own their homes. * Their employment types.

```
In [116]: sns.distplot(train_new['Age'], hist_kws=dict(edgecolor="k", linewidth=2))
          plt.ylabel("Density estimate", fontsize = 16)
          plt.xlabel('Age', fontsize = 16)
          plt.title("Distribution of applicants - AGE", fontsize = 22)
          plt.show()
```



The bulk of applicants seem to be under 40, with 25 – 38 accounting for the largest segment. Now let's see if Age has any impact on WasTheLoanApproved.

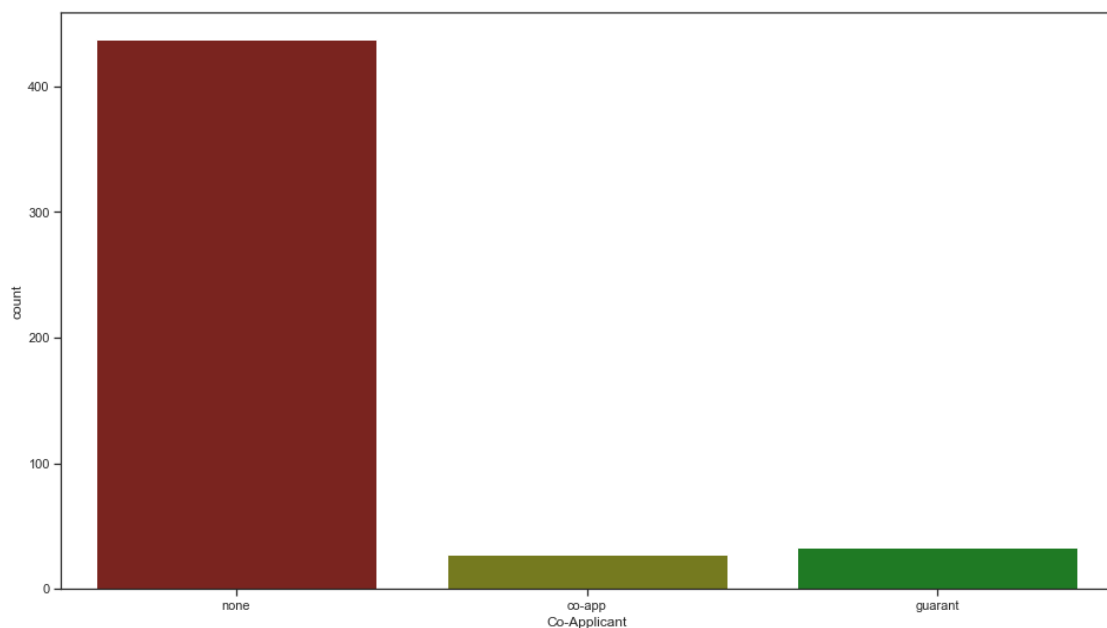
```
In [117]: g = sns.catplot(x = "WasTheLoanApproved", y = "Age", kind = "swarm", data = train_new)
g.fig.set_size_inches(16, 9)
```



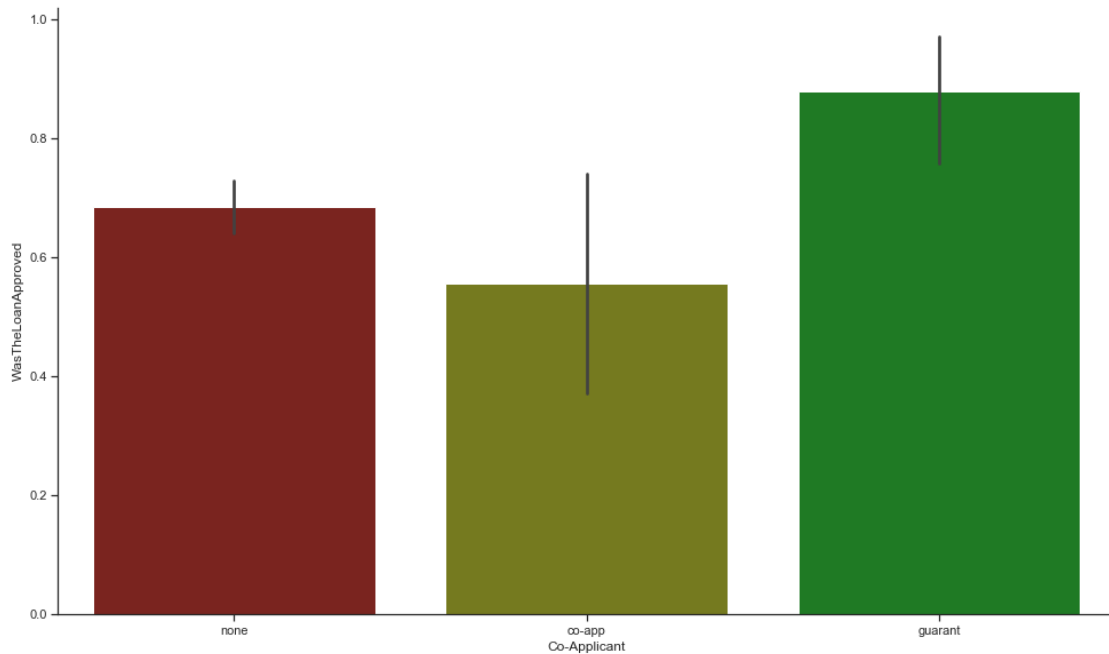
Since both groups are visually similar, it is difficult to say if Age played a decision in the loan approval process from this graphic alone. Let's explore Co-Applicant and see if applying with a Co-Applicant improves an applicant's odds of approval.

```
In [118]: sns.countplot(train_new["Co-Applicant"])
```

```
Out[118]: <matplotlib.axes._subplots.AxesSubplot at 0x168a2fb3470>
```

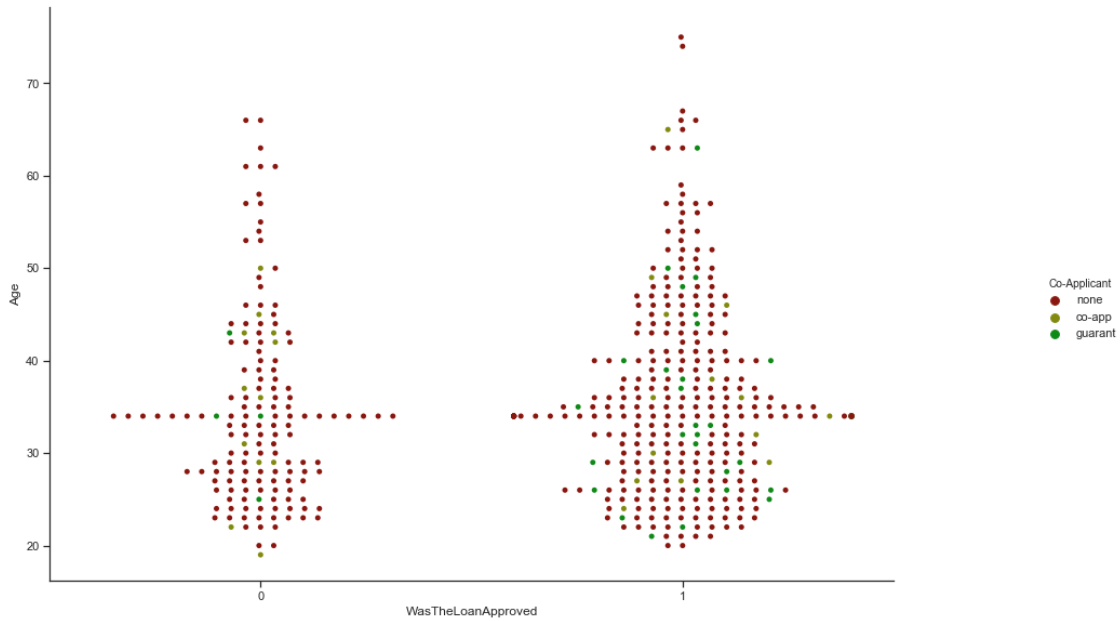


```
In [119]: g = sns.catplot(x = "Co-Applicant", y = "WasTheLoanApproved", kind="bar", data = tra
          g.fig.set_size_inches(16, 9)
```



As we can see, the vast majority of applicants apply for a loan by themselves. However, applying with a guarantor improves your chances of approval. Let us pair this with Age. This will help add more weight to our assumption.

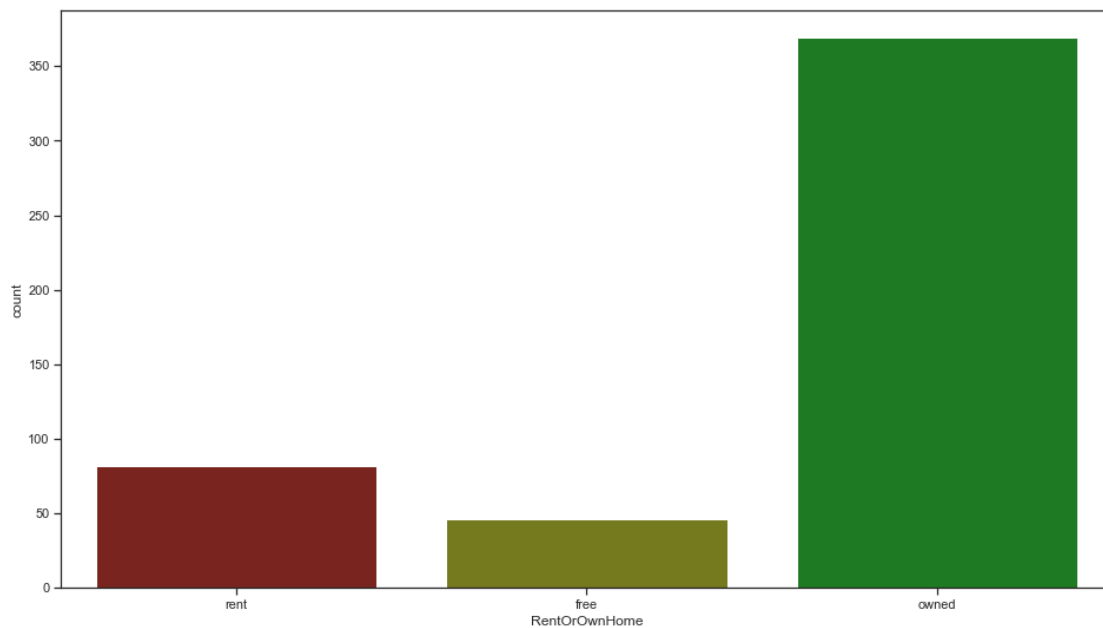
```
In [120]: g = sns.catplot(x = "WasTheLoanApproved", y = "Age", hue = "Co-Applicant", kind = "s
          g.fig.set_size_inches(16, 9)
```



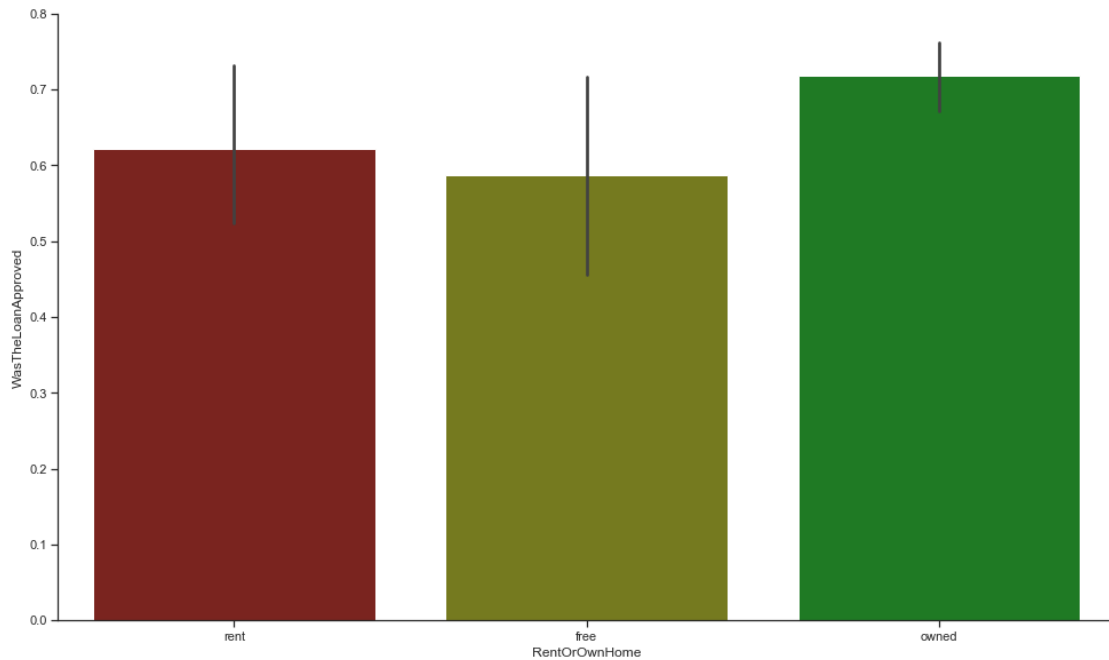
As we can see, the number of guarant points for rejected applicants is much lower than those for approved applicants. Having a guarantor as a Co-Applicant certainly seems to have an impact on our target variable. Let us now explore whether our applicants rent/own their homes and what role that feature may play in our analysis.

```
In [121]: sns.countplot(train_new["RentOrOwnHome"])
```

```
Out[121]: <matplotlib.axes._subplots.AxesSubplot at 0x168a7443438>
```



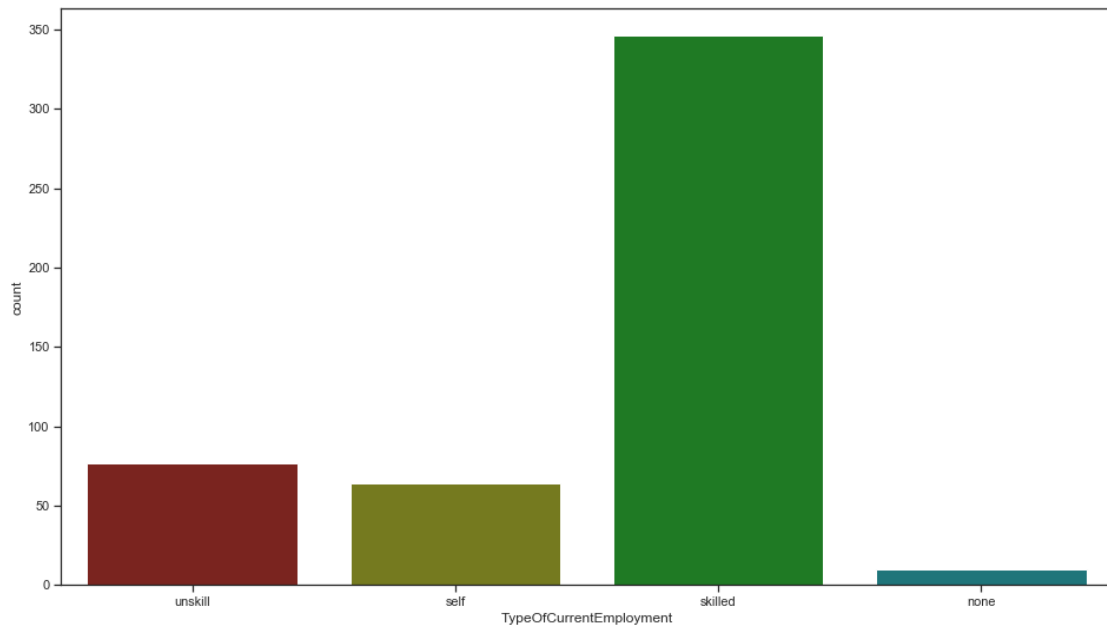
```
In [122]: g = sns.catplot(x = "RentOrOwnHome", y = "WasTheLoanApproved", kind="bar", data = tra
          g.fig.set_size_inches(16, 9)
```



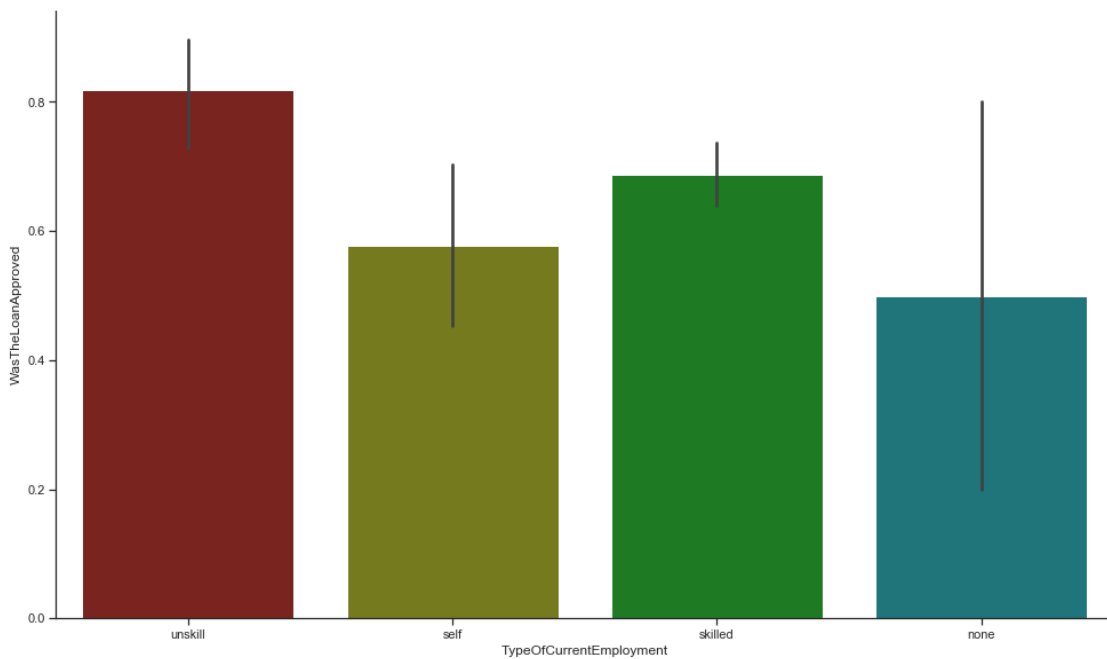
Owning a home seems to have an impact on our target variable. Let us now explore what kind of employment types our applicants have and what role that feature may play in our analysis.

```
In [123]: sns.countplot(train_new["TypeOfCurrentEmployment"])
```

```
Out[123]: <matplotlib.axes._subplots.AxesSubplot at 0x168a8b43cf8>
```

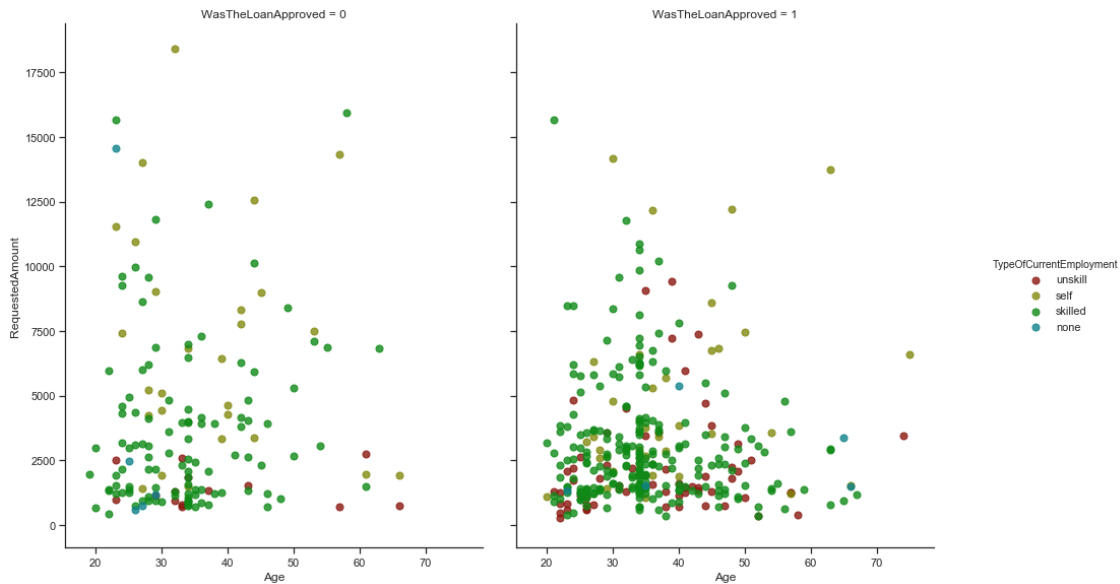


```
In [124]: g = sns.catplot(x = "TypeOfCurrentEmployment", y = "WasTheLoanApproved", kind="bar",  
g.fig.set_size_inches(16, 9)
```



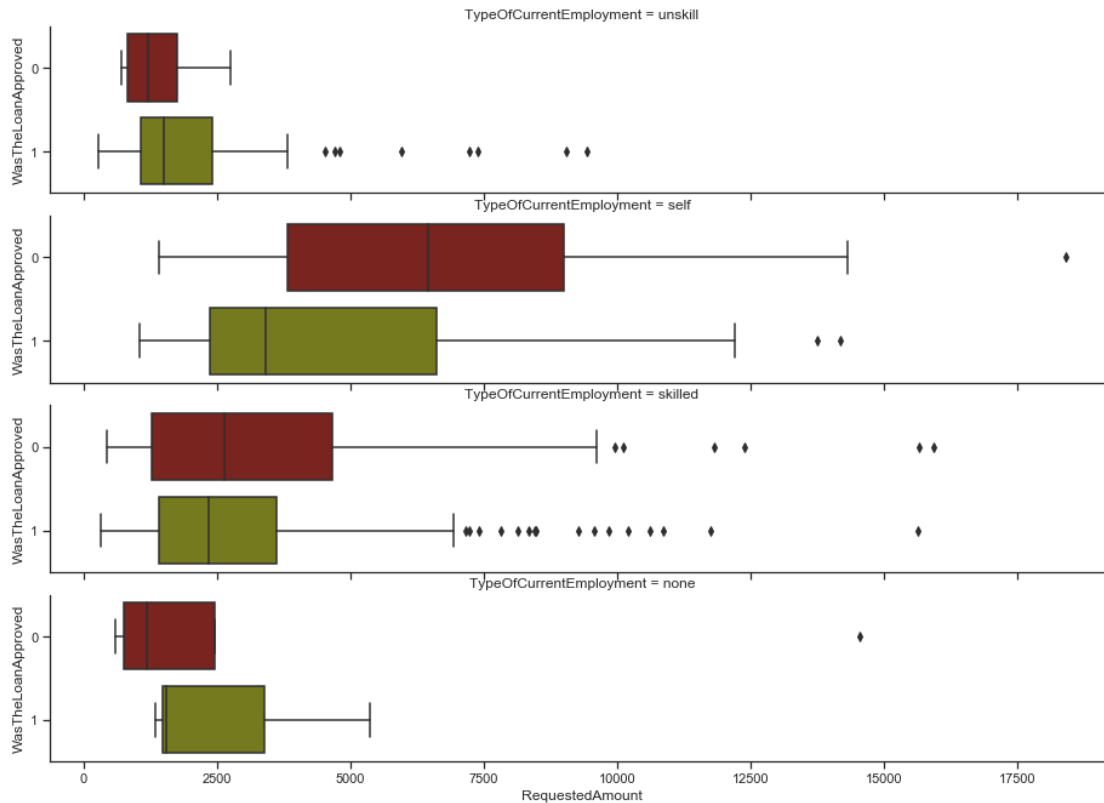
Let us explore this further by pairing this data with Age and RequestedAmount.

```
In [125]: g = sns.lmplot(data = train_new, x = "Age", y = "RequestedAmount",
                        hue = "TypeOfCurrentEmployment", col="WasTheLoanApproved",
                        scatter_kws = {'s':50}, fit_reg=False)
g.fig.set_size_inches(16, 9)
```



This plot doesn't tell us much apart from - * Self employed applicants who get rejected apply for larger amounts than other employment type applicants. When self-employed applicants apply for lesser amounts (< \$7500), they have a greater chance of approval. * Younger applicants, irrespective of employment type have a good likelihood of being rejected if they apply for large amounts (> \$7500)

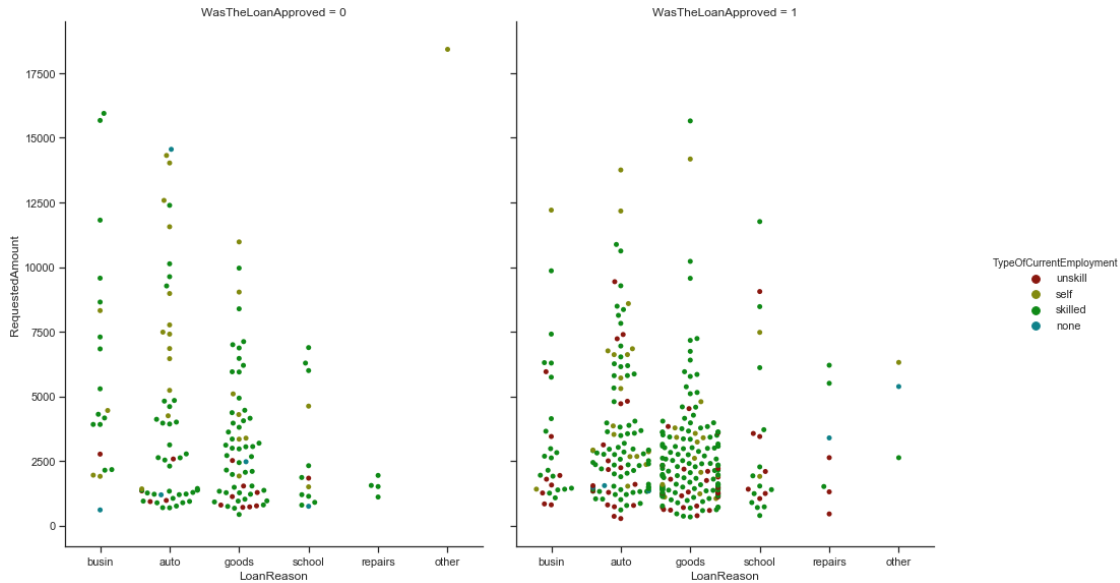
```
In [126]: g = sns.catplot(x="RequestedAmount", y="WasTheLoanApproved", row="TypeOfCurrentEmployment",
                        kind="box", orient="h", height = 4, aspect = 0.9 , data = train_new)
g.fig.set_size_inches(16, 9)
```



The above plot offers some interesting insights - * Applicants with unskilled employment tend to apply for loans with amounts below \$5000, a low amount seems to improve their chances of approval. The mean amount for which they do get approved is the lowest « \$2500. * Applicants who are currently unemployed tend to apply for loans with amounts up to \$15000. However, probably due to their unemployed status, almost all loan applications with the amount > \$5000 seem to lead to a rejection. The mean amount for which they do get approved is the 2nd lowest < \$2500. * The highest mean loan amount that is approved is for the self employed applicants. They also have the highest mean loan amount that gets rejected, leading to the assumption that they tend to apply for loans with higher amounts than other categories.

Let us look at some of the **reasons** these categories of applicants, apply for loans by adding LoanReason to the mix.

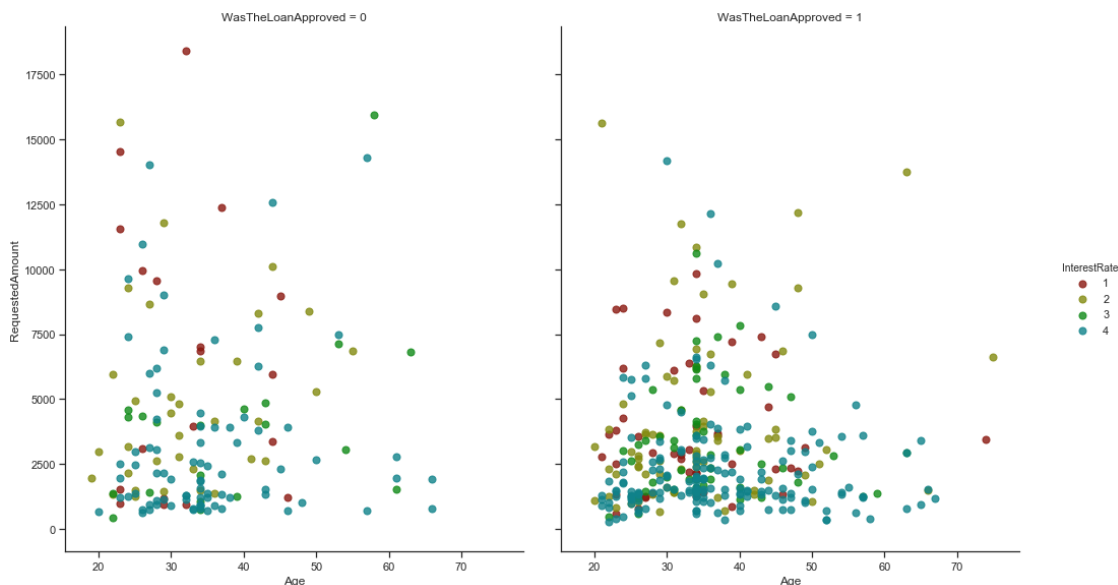
```
In [127]: g = sns.catplot(x="LoanReason", y="RequestedAmount", hue="TypeOfCurrentEmployment",
                        col="WasTheLoanApproved",
                        kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```

Once again, this is not very clear. Here are some assumptions we can draw from these visuals. * Unskilled applicants tend to form a significant proportion of people applying for school or business loans. They apply for small amounts and tend to get approved. * When self-employed applicants apply for auto or goods loans with amounts > \$5000, they tend to get rejected.

Let us shift our focus to the finances of our applicants and see if we can isolate any patterns in those features. Let us first look at InterestRate.

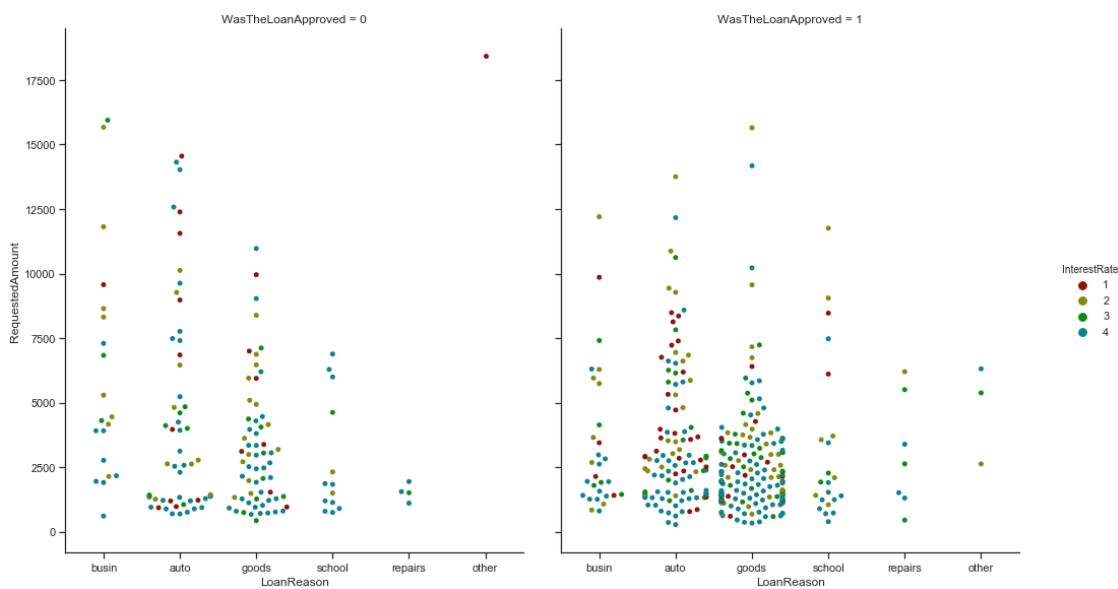
```
In [128]: g = sns.lmplot(data = train_new, x = "Age", y = "RequestedAmount",
                        hue = "InterestRate", col="WasTheLoanApproved",
                        scatter_kws = {'s':50}, fit_reg=False)
g.fig.set_size_inches(16, 9)
```



Unsurprisingly, the bank seems to approve **high interest, low amount** loans and tends to reject **low interest, high amount** loans. This seems to be a way to mitigate overall risk as a large number of **high interest, low amount** loans - * Bring in more revenue for the bank * Are less risky.

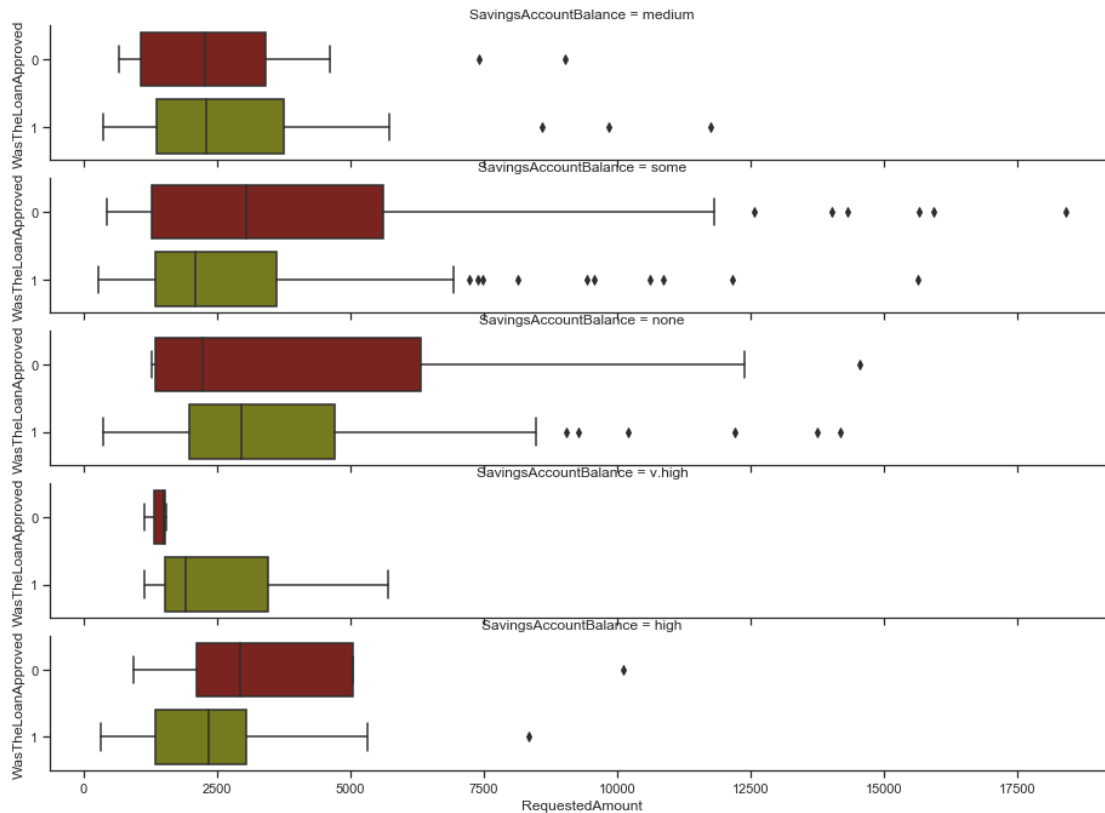
We can also see that almost all applications where the RequestedAmount > \$7500 and the InterestRate == > 1% were rejected. Probably because a few **low interest, high amount** loans do not bring in enough revenue to justify the risk.

```
In [129]: g = sns.catplot(x="LoanReason", y="RequestedAmount", hue="InterestRate",
                        col="WasTheLoanApproved",
                        kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```



Let us explore SavingsAccountBalance, in my opinion it is a stronger indicator than CheckingAccountBalance because - * CheckingAccountBalance will have a higher degree of **churn** than SavingsAccountBalance, i.e, Savings account balances may not decrease, but grow over time, whereas CheckingAccountBalance may follow monthly cyclical patterns due to bills, expenses etc. * Since this is categorical data, the likelihood of savings account balances jumping from category to category is lower than the same for checking account balances.

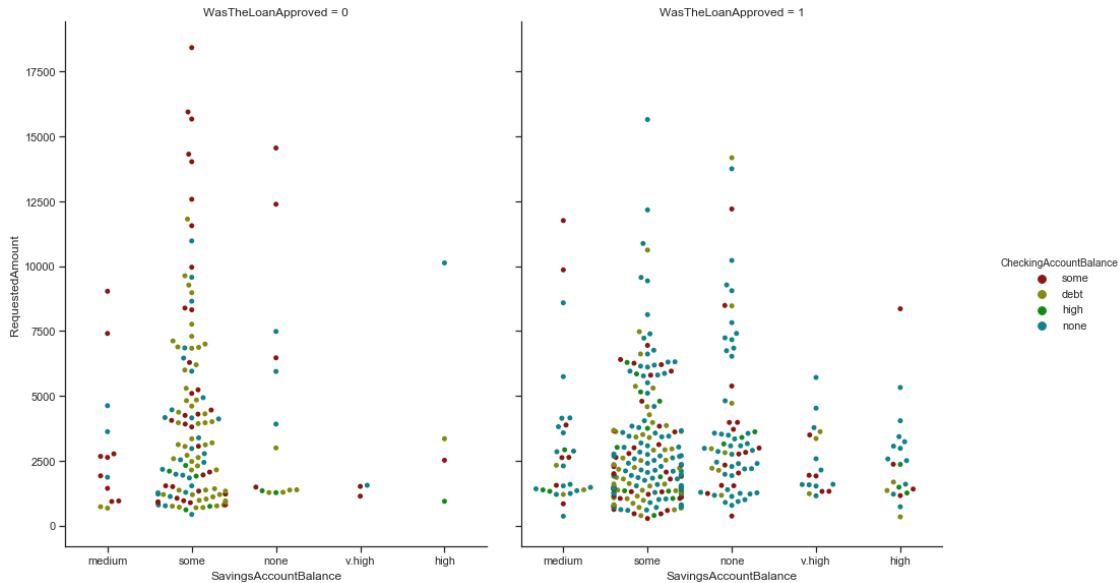
```
In [130]: g = sns.catplot(x="RequestedAmount", y="WasTheLoanApproved", row="SavingsAccountBalance",
                        kind="box", orient="h", height = 4, aspect = 0.9 , data = train_new)
g.fig.set_size_inches(16, 9)
```



The above graphic is very helpful and we can make the following assumptions - * Applicants with some or none or medium savings tend to apply for loans of all amounts but get rejected for loans of higher amounts. * Applicants with high or v. high savings tend to apply for loans of small amounts, however surprisingly on average, they tend to get rejected for smaller loans.

From these observations, we can surmise that SavingsAccountBalance plays a clear, important role in the loan approval process. Let us now add CheckingAccountBalance to the mix and see how that changes our picture.

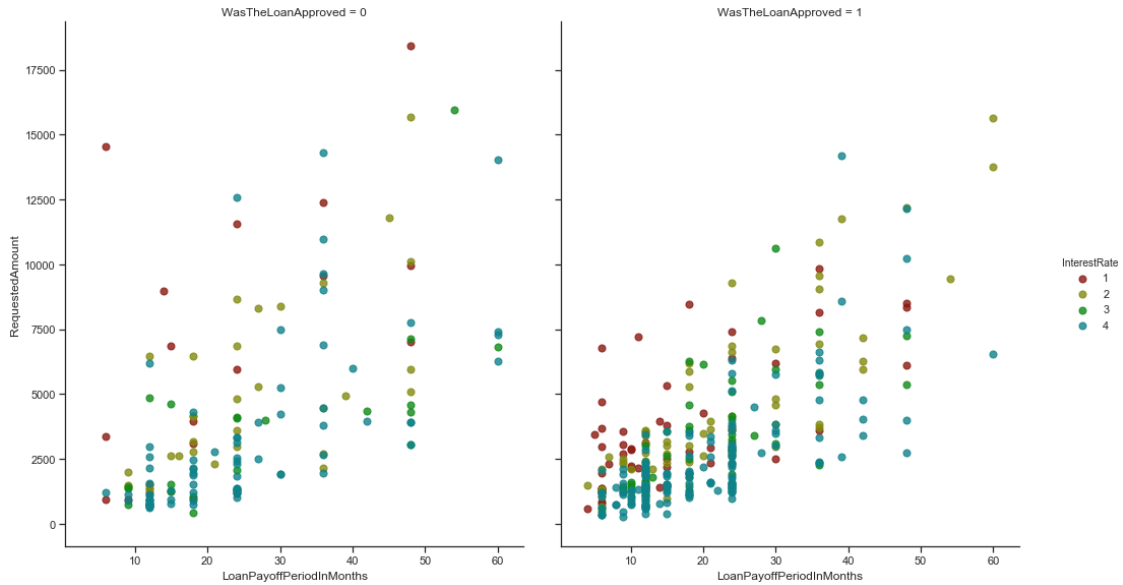
```
In [131]: g = sns.catplot(x = "SavingsAccountBalance", y="RequestedAmount", hue="CheckingAccountBalance",
                        col="WasTheLoanApproved",
                        kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```



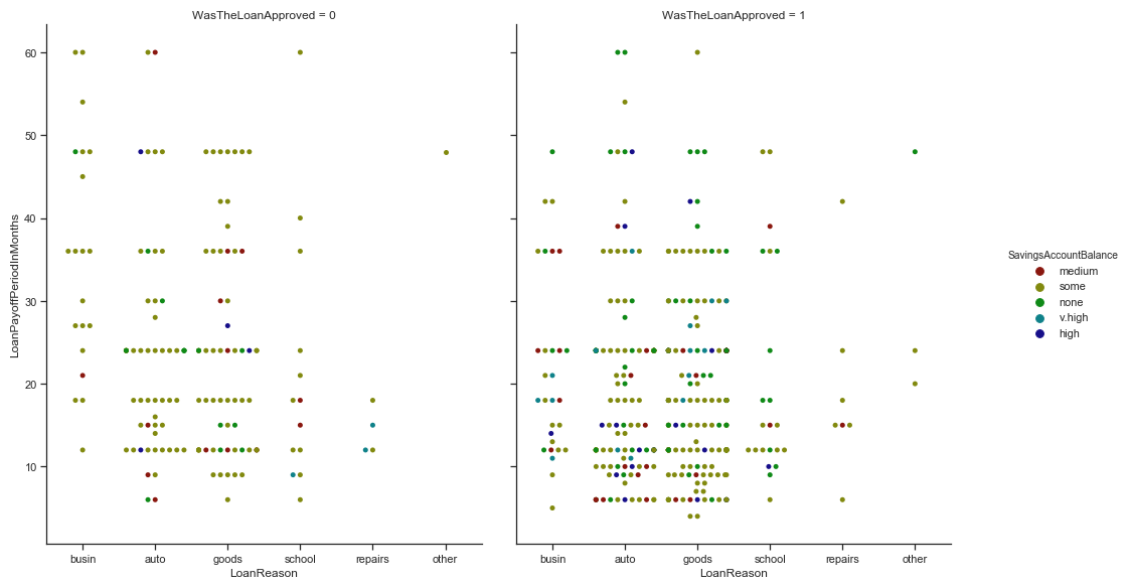
We can observe - * Barring some cases, having debt in CheckingAccountBalance seems to lead to a rejection of the loan application. * Applicants with no checking account balance tend to borrow low amounts and do get approved. It is similar for people with people with high checking account balance.

Let us now explore the relationships between LoanPayoffPeriodInMonths and other features. In the plot below we explore how LoanPayoffPeriodInMonths relates to RequestedAmount and how that may impact approval. We can observe that loans for higher time periods have higher rates. Also, short time period loans, with low amounts tend to get approved while shorter loans with low interest rates for higher values tend to get rejected.

```
In [132]: g = sns.lmplot(data = train_new, x = "LoanPayoffPeriodInMonths", y = "RequestedAmount",
                        hue = "InterestRate", col="WasTheLoanApproved",
                        scatter_kws = {'s':50}, fit_reg=False)
g.fig.set_size_inches(16, 9)
```



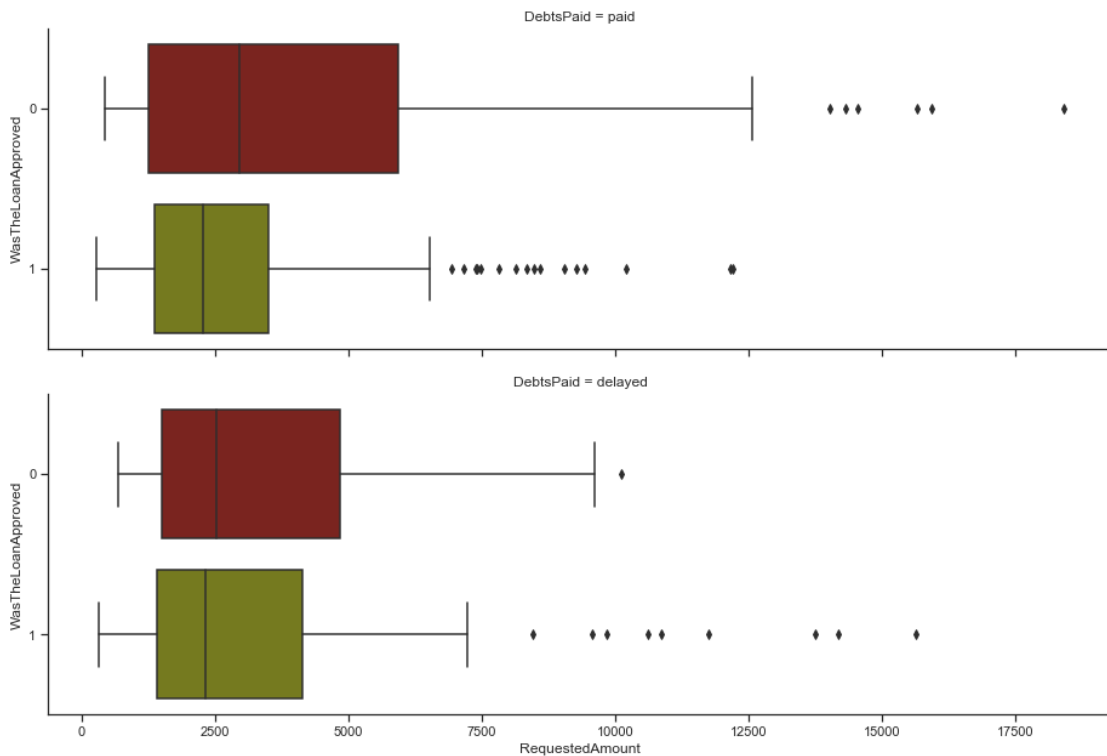
```
In [133]: g = sns.catplot(x = "LoanReason", y="LoanPayoffPeriodInMonths", hue="SavingsAccountBalance",
                        col="WasTheLoanApproved",
                        kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```



The above graphic helps us understand that - * Long term auto loans tend to get rejected probably because vehicles depreciate in value very rapidly. * Applicants with very high savings tend to get approved for long term loans.

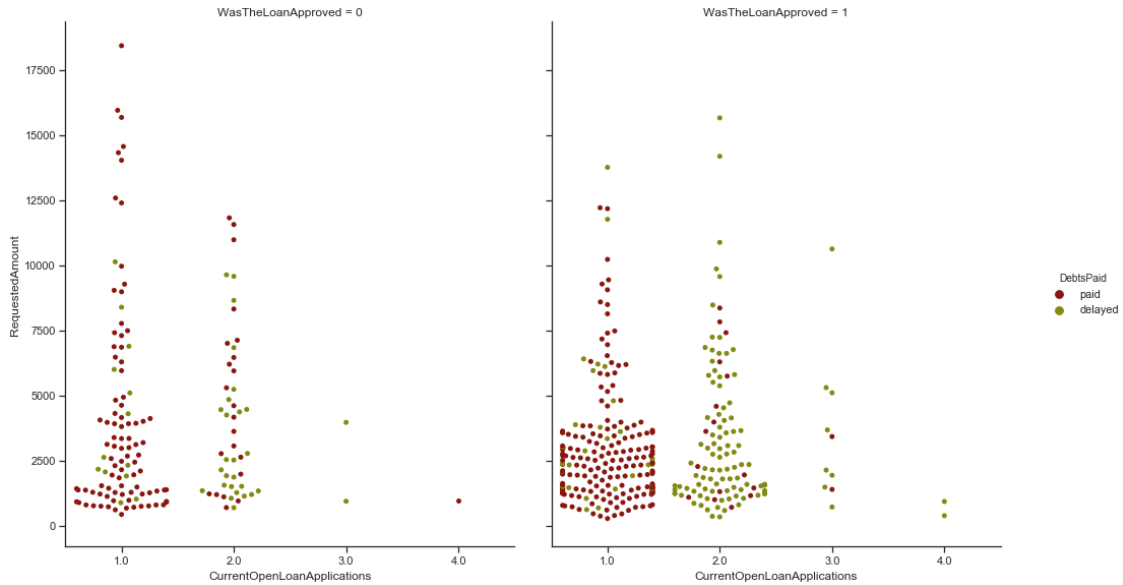
Finally, let us explore DebtsPaid, CurrentOpenLoanApplications and see how they relate to the other features.

```
In [134]: g = sns.catplot(x="RequestedAmount", y="WasTheLoanApproved", row="DebtsPaid",
                        kind="box", orient="h", height = 4, aspect = 0.9 , data = train_new)
g.fig.set_size_inches(16, 9)
```



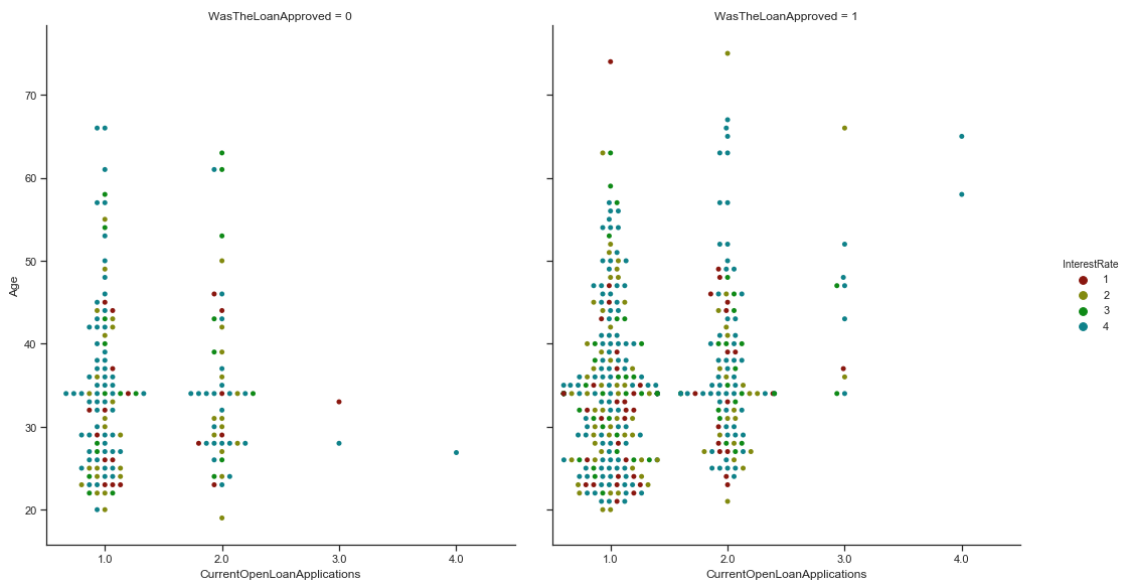
Surprisingly, getting delayed on paying back debt doesn't seem to have the effect we're looking for. It seems that on average, applicants with a delay in paying back their debts have a similar loan amount approved as those who pay their debts on time.

```
In [135]: g = sns.catplot(x = "CurrentOpenLoanApplications", y="RequestedAmount", hue="DebtsPaid",
                        col="WasTheLoanApproved",
                        kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```



Having multiple open loan applications along with having delayed payments also doesn't seem to impact loan approval as there are plenty of applicants who have been approved inspite of being delayed and having multiple open loans. Unfortunately, applicants with only a single loan, who have paid off their debts seem to have a high probability of getting rejected when they apply for loans with high amounts.

```
In [136]: g = sns.catplot(x = "CurrentOpenLoanApplications", y="Age", hue="InterestRate",
    col="WasTheLoanApproved",
    kind="swarm", data=train_new)
g.fig.set_size_inches(16, 9)
```



From this graphic we can observe - * Having 2 or more open loan applications leads to a loan approval with higher interest rates. * Applicants older than 30 tend to have ≥ 3 open loan applications. They usually get approved but with higher interest rates.

1.2 Building our Classifiers

```
In [64]: train_new.dtypes
        # train_new
```

```
Out [64]: LoanPayoffPeriodInMonths      int64
LoanReason                             object
RequestedAmount                        int64
InterestRate                          int64
Co-Applicant                          object
WasTheLoanApproved                    int64
YearsInCurrentResidence                float64
Age                                    float64
NumberOfDependantsIncludingSelf          float64
CurrentOpenLoanApplications           float64
YearsAtCurrentEmployer                object
RentOrOwnHome                         object
TypeOfCurrentEmployment               object
CheckingAccountBalance                object
DebtsPaid                             object
SavingsAccountBalance                 object
dtype: object
```

We use **one hot encoding** to convert categorical variables to numerically encoded values for our model prediction

```
In [65]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
obj_dtypes = [i for i in train_new.select_dtypes(include = np.object).columns]
for i in obj_dtypes:
    train_new[i] = le.fit_transform(train_new[i])
    test_new[i] = le.fit_transform(test_new[i])
    blind_new[i] = le.fit_transform(blind_new[i])
train_new = pd.get_dummies(data=train_new,columns=obj_dtypes)
test_new = pd.get_dummies(data=test_new,columns=obj_dtypes)
blind_new = pd.get_dummies(data=blind_new,columns=obj_dtypes)
```

```
In [66]: train_new.dtypes
```

```
Out [66]: LoanPayoffPeriodInMonths      int64
RequestedAmount                        int64
InterestRate                          int64
WasTheLoanApproved                    int64
YearsInCurrentResidence                float64
```


Age	float64
NumberOfDependantsIncludingSelf	float64
CurrentOpenLoanApplications	float64
LoanReason_0	uint8
LoanReason_1	uint8
LoanReason_2	uint8
LoanReason_3	uint8
LoanReason_4	uint8
LoanReason_5	uint8
Co-Applicant_0	uint8
Co-Applicant_1	uint8
Co-Applicant_2	uint8
YearsAtCurrentEmployer_0	uint8
YearsAtCurrentEmployer_1	uint8
YearsAtCurrentEmployer_2	uint8
YearsAtCurrentEmployer_3	uint8
YearsAtCurrentEmployer_4	uint8
RentOrOwnHome_0	uint8
RentOrOwnHome_1	uint8
RentOrOwnHome_2	uint8
TypeOfCurrentEmployment_0	uint8
TypeOfCurrentEmployment_1	uint8
TypeOfCurrentEmployment_2	uint8
TypeOfCurrentEmployment_3	uint8
CheckingAccountBalance_0	uint8
CheckingAccountBalance_1	uint8
CheckingAccountBalance_2	uint8
CheckingAccountBalance_3	uint8
DebtsPaid_0	uint8
DebtsPaid_1	uint8
SavingsAccountBalance_0	uint8
SavingsAccountBalance_1	uint8
SavingsAccountBalance_2	uint8
SavingsAccountBalance_3	uint8
SavingsAccountBalance_4	uint8
dtype:	object

In [67]: test_new.dtypes

Out [67]:	LoanPayoffPeriodInMonths	int64
	RequestedAmount	int64
	InterestRate	int64
	WasTheLoanApproved	int64
	YearsInCurrentResidence	float64
	Age	float64
	NumberOfDependantsIncludingSelf	float64
	CurrentOpenLoanApplications	float64
	LoanReason_0	uint8

LoanReason_1	uint8
LoanReason_2	uint8
LoanReason_3	uint8
LoanReason_4	uint8
LoanReason_5	uint8
Co-Applicant_0	uint8
Co-Applicant_1	uint8
Co-Applicant_2	uint8
YearsAtCurrentEmployer_0	uint8
YearsAtCurrentEmployer_1	uint8
YearsAtCurrentEmployer_2	uint8
YearsAtCurrentEmployer_3	uint8
YearsAtCurrentEmployer_4	uint8
RentOrOwnHome_0	uint8
RentOrOwnHome_1	uint8
RentOrOwnHome_2	uint8
TypeOfCurrentEmployment_0	uint8
TypeOfCurrentEmployment_1	uint8
TypeOfCurrentEmployment_2	uint8
TypeOfCurrentEmployment_3	uint8
CheckingAccountBalance_0	uint8
CheckingAccountBalance_1	uint8
CheckingAccountBalance_2	uint8
CheckingAccountBalance_3	uint8
DebtsPaid_0	uint8
DebtsPaid_1	uint8
SavingsAccountBalance_0	uint8
SavingsAccountBalance_1	uint8
SavingsAccountBalance_2	uint8
SavingsAccountBalance_3	uint8
SavingsAccountBalance_4	uint8
dtype: object	

In [68]: blind_new.dtypes

Out [68]:	LoanPayoffPeriodInMonths	int64
	RequestedAmount	int64
	InterestRate	int64
	WasTheLoanApproved	int64
	YearsInCurrentResidence	float64
	Age	float64
	NumberOfDependantsIncludingSelf	float64
	CurrentOpenLoanApplications	float64
	LoanReason_0	uint8
	LoanReason_1	uint8
	LoanReason_2	uint8
	LoanReason_3	uint8
	LoanReason_4	uint8

LoanReason_5	uint8
Co-Applicant_0	uint8
Co-Applicant_1	uint8
Co-Applicant_2	uint8
YearsAtCurrentEmployer_0	uint8
YearsAtCurrentEmployer_1	uint8
YearsAtCurrentEmployer_2	uint8
YearsAtCurrentEmployer_3	uint8
YearsAtCurrentEmployer_4	uint8
RentOrOwnHome_0	uint8
RentOrOwnHome_1	uint8
RentOrOwnHome_2	uint8
TypeOfCurrentEmployment_0	uint8
TypeOfCurrentEmployment_1	uint8
TypeOfCurrentEmployment_2	uint8
TypeOfCurrentEmployment_3	uint8
CheckingAccountBalance_0	uint8
CheckingAccountBalance_1	uint8
CheckingAccountBalance_2	uint8
CheckingAccountBalance_3	uint8
DebtsPaid_0	uint8
DebtsPaid_1	uint8
SavingsAccountBalance_0	uint8
SavingsAccountBalance_1	uint8
SavingsAccountBalance_2	uint8
SavingsAccountBalance_3	uint8
SavingsAccountBalance_4	uint8
dtype:	object

As we can see, one hot encoding, expands our data set by applying coded values to categorical variables. We now separate the predictor variables (X) and target variable (Y) into separate dataframe for our train, test and blind data sets.

```
In [69]: x_train = train_new[[i for i in train_new.columns if i not in ['WasTheLoanApproved']]]
        y_train = train_new.WasTheLoanApproved

        x_test = test_new[[i for i in test_new.columns if i not in ['WasTheLoanApproved']]]
        y_test = test_new.WasTheLoanApproved

        x_blind = blind_new[[i for i in blind_new.columns if i not in ['WasTheLoanApproved']]]
        y_blind = blind_new.WasTheLoanApproved
```

In the snippet below, I write a model function that takes as input the classification algorithm, X/Y train/test/blind data sets. The function - 1. Fits the model to the training dataset. 2. Makes predictions for the test dataset. 3. Makes predictions for the blind dataset. 4. Calculates quality metrics * Accuracy score * Recall score * Classification report containing a detailed summary. 5. Plots the AUC and ROC Curve. 6. Writes the test predictions to a file. 7. Writes the blind predictions to a file.

```

In [70]: from sklearn.metrics import confusion_matrix, accuracy_score, recall_score, roc_auc_score

def model(algorithm, dtrain_X, dtrain_Y, dtest_X, dtest_Y, blind_X, blind_Y, algo = 0):
    cols = dtrain_X.columns
    algorithm.fit(dtrain_X[cols], dtrain_Y)
    predictions = algorithm.predict(dtest_X[cols]) # Predictions for the test dataset
    blindPredictions = algorithm.predict(blind_X[cols]) # Predictions for the blind dataset
    if(iterative == False):
        print (algorithm)

        print ("Accuracy score : ", round(accuracy_score(dtest_Y, predictions), 3))
        print ("Recall score : ", round(recall_score(dtest_Y, predictions), 3))
        print ("classification report :\n", classification_report(dtest_Y, predictions))

        fig = plt.figure(figsize=(16, 9))
        ax = fig.add_subplot(111)
        prediction_probabilities = algorithm.predict_proba(dtest_X[cols])[:,1]
        fpr, tpr, thresholds = roc_curve(dtest_Y, prediction_probabilities)
        ax.plot(fpr, tpr, label = ["Area under curve : ", round(auc(fpr, tpr), 3)], linewidth=4, linestyle="dashed")
        ax.plot([0,1],[0,1],linewidth = 4, linestyle = "dashed")
        plt.legend(loc = "best")
        plt.title("ROC-CURVE & AREA UNDER CURVE")

        fname = ""
        if(algo == 1):
            fname = "logit"
        elif(algo == 2):
            fname = "randomForest"

        s0 = dtest_Y.reset_index(level=0)
        s1 = pd.Series(predictions, name = 'predicted')
        pd.concat([s0, s1], axis = 1).to_csv(fname + "modelPred.csv", sep=',', index = False)

        s2 = blind_Y.reset_index(level=0)
        s3 = pd.Series(blindPredictions, name = 'predicted')
        pd.concat([s2, s3], axis = 1).to_csv(fname + "blindPred.csv", sep=',', index = False)

    return classification_report(dtest_Y, predictions, output_dict = True)

```

1.2.1 Logistic Regression

```

In [71]: from sklearn.linear_model import LogisticRegression
logit = LogisticRegression()
logitModel = model(logit, x_train, y_train, x_test, y_test, x_blind, y_blind, algo = 1)

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=None, solver='warn',

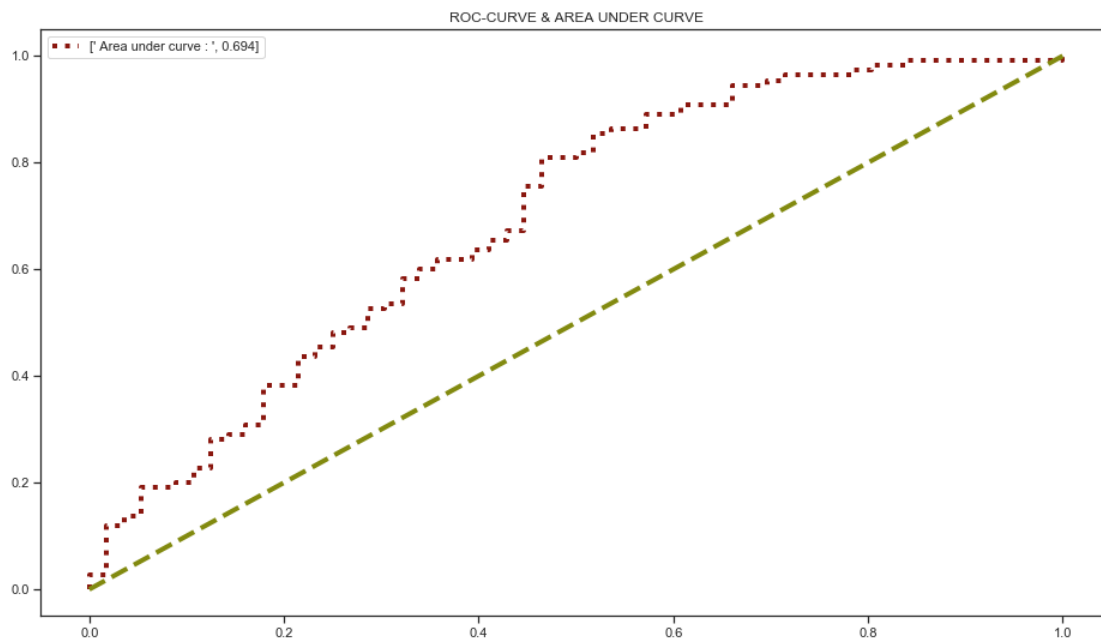
```

```

        tol=0.0001, verbose=0, warm_start=False)
Accuracy score : 0.735
Recall score   : 0.945
classification report :

```

	precision	recall	f1-score	support
0	0.75	0.32	0.45	56
1	0.73	0.95	0.83	110
micro avg	0.73	0.73	0.73	166
macro avg	0.74	0.63	0.64	166
weighted avg	0.74	0.73	0.70	166



These are the results of the **Logistic Regression** classifier. Our accuracy is near 80%. However, as our data has some imbalance, better metrics to judge our model are **AUC** (area under the ROC curve), **Weighted-Avg. Precision**, **Weighted-Avg. Recall** and **Weighted-Avg. f1-score**. In the next snippet of code, I run this model on our data 10,000 times. This isn't necessary to do, but it gives us strong **average** values of what precision, recall and f1-score look like over multiple iterations.

```

In [251]: start = time.time()
          precision = 0
          recall = 0
          f1 = 0
          runs = 10000
          for i in range(runs):
              logitModel = model(logit, x_train, y_train, x_test, y_test, x_blind, y_blind, it

```

```

        precision += logitModel['weighted avg']['precision']
        recall += logitModel['weighted avg']['recall']
        f1 += logitModel['weighted avg']['f1-score']
    stop = time.time()
    logitpAvg = round(precision/runs, 4)
    logitrAvg = round(recall/runs, 4)
    logitfAvg = round(f1/runs, 4)

In [252]: print("Precision Average :", logitpAvg)
          print("Recall Average :", logitrAvg)
          print("f1-score Average :", logitfAvg)
          print("Time Elapsed :", round(stop - start, 3), "seconds.")

Precision Average : 0.7262
Recall Average : 0.741
f1-score Average : 0.7312
Time Elapsed : 189.388 seconds.

```

These values are averaged out over 10,000 iterations and are indicative of the precision, recall values we would see if we applied our model to the blind loan dataset. These numbers are indicative of the quality of this model and one way to judge if this model is any good is to compare it against a baseline. We will implement a Monte-Carlo random classifier later and see if this model beats our baseline.

1.2.2 Random Forest Classifier

```

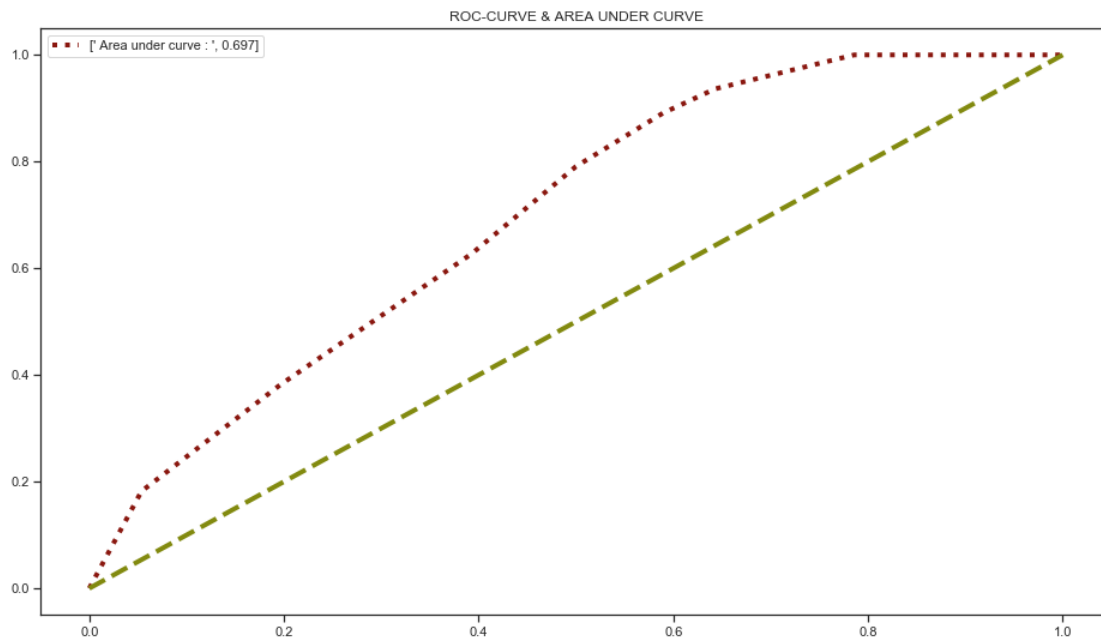
In [72]: from sklearn.ensemble import RandomForestClassifier
         rfc = RandomForestClassifier()
         rfcModel = model(rfc, x_train, y_train, x_test, y_test, x_blind, y_blind, algo = 2)

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)

Accuracy score : 0.729
Recall score   : 0.891
classification report :

```

	precision	recall	f1-score	support
0	0.66	0.41	0.51	56
1	0.75	0.89	0.81	110
micro avg	0.73	0.73	0.73	166
macro avg	0.70	0.65	0.66	166
weighted avg	0.72	0.73	0.71	166



These are the results of the **Random Forest** classifier. Our accuracy is near 75%. However, as our data has some imbalance, better metrics to judge our model are **AUC** (area under the ROC curve), **Weighted-Avg. Precision**, **Weighted-Avg. Recall** and **Weighted-Avg. f1-score**. In the next snippet of code, I run this model on our data 10,000 times. Once again this isn't necessary to do, but it gives us strong **average** values of what precision, recall and f1-score look like over multiple iterations.

```
In [261]: start = time.time()
precision = 0
recall = 0
f1 = 0
runs = 10000
for i in range(runs):
    rfcModel = model(rfc, x_train, y_train, x_test, y_test, x_blind, y_blind, iterat
    precision += rfcModel['weighted avg']['precision']
    recall += rfcModel['weighted avg']['recall']
    f1 += rfcModel['weighted avg']['f1-score']
stop = time.time()
rfcpAvg = round(precision/runs, 4)
rfcrAvg = round(recall/runs, 4)
rfcfAvg = round(f1/runs, 4)

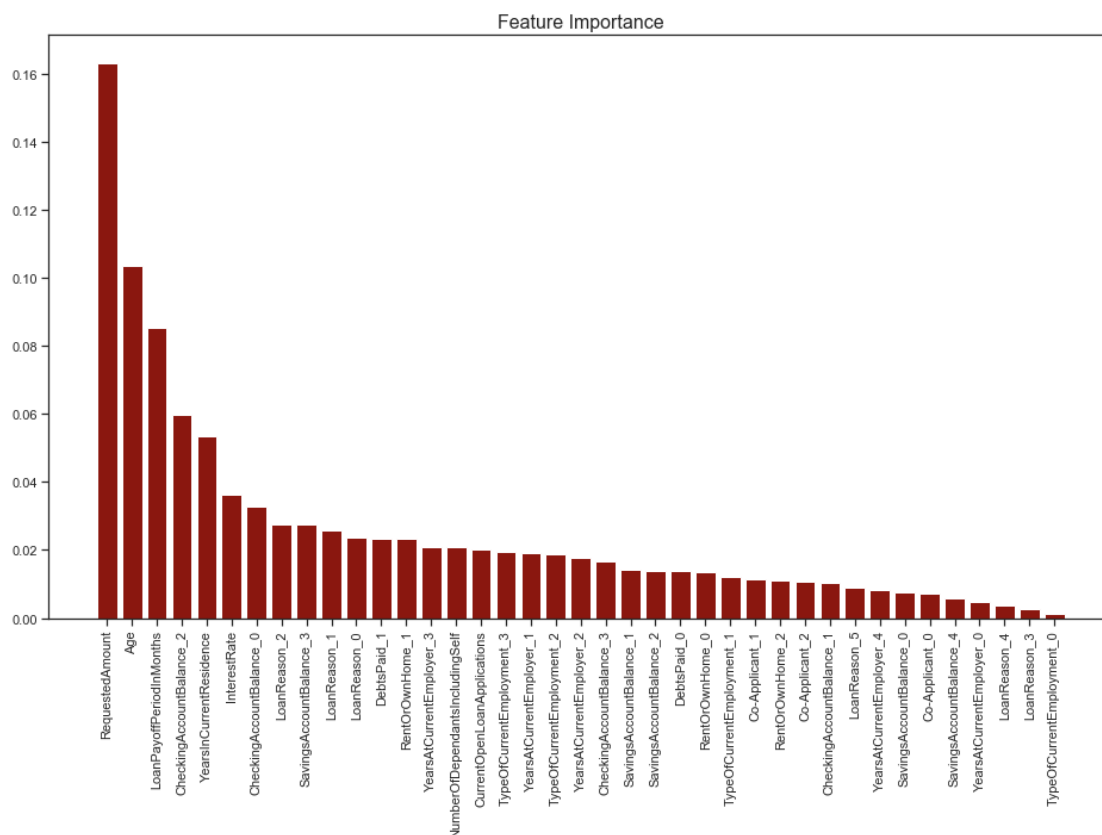
In [262]: print("Precision Average :", rfcpAvg)
print("Recall Average :", rfcrAvg)
print("f1-score Average :", rfcfAvg)
print("Time Elapsed :", round(stop - start, 3), "seconds.")
```

```
Precision Average : 0.7036
Recall Average : 0.6931
f1-score Average : 0.6973
Time Elapsed : 392.723 seconds.
```

Using a Random Forest Classifier, we can visualize the importance of the features that we have used in building our model.

```
In [73]: # Plot features importance
importances = rfc.feature_importances_
indices = np.argsort(rfc.feature_importances_)[::-1]
plt.bar(range(1, len(importances)+1), importances[indices], align="center")
plt.xticks(range(1, len(importances)+1),
           train_new.columns[train_new.columns != "WasTheLoanApproved"][indices],
           rotation=90)
plt.title("Feature Importance", {"fontsize": 16})
```

```
Out[73]: Text(0.5, 1.0, 'Feature Importance')
```



The plot above tells us which features contributed were the most important for our Random Forest Model. RequestedAmount, LoanPeriod and Age played the most important roles. The top 5 most important features are all numerical in type.

A plot like this is helpful in many ways, namely - * It helps us judge if the assumptions we made at the beginning were correct. * It tells us which features we can remove, if they do not contribute value to our model. * It can help us refine our approach and our model.

In the plot above, we can see that almost all features have some degree of importance in the overall model. However, if we did have a case where we have many features, we could use a plot like this to choose the most important features and act accordingly.

1.2.3 Monte-Carlo Random Prediction baseline

One way to check if our classifiers are any good, is to answer the question, “**Are our results better than random guesses?**” To answer this question, we randomly assign a loan approval/denial to the customers in our test set and compare against the actual results.

```
In [264]: import random
          randomGuess = [random.randint(0, 1) for _ in range(x_test.shape[0])] #A random loan
          print(classification_report(y_test, randomGuess))
```

	precision	recall	f1-score	support
0	0.26	0.40	0.31	45
1	0.72	0.57	0.64	121
micro avg	0.52	0.52	0.52	166
macro avg	0.49	0.49	0.47	166
weighted avg	0.59	0.52	0.55	166

Looking at the weighted averages, we can confirm that the values are lower than those given by our logistic regression and random forest classifiers. Now, by running this 5,000 or 10,000 times, we can get a better picture of what a truly **random** guess implies and whether our classifiers beat such a value and are therefore better than a random guess.

```
In [265]: start = time.time()
          precision = 0
          recall = 0
          f1 = 0
          runs = 10000
          for i in range(runs):
              randomGuess = [random.randint(0, 1) for _ in range(x_test.shape[0])]
              report = classification_report(y_test, randomGuess, output_dict = True)
              precision += report['weighted avg']['precision']
              recall += report['weighted avg']['recall']
              f1 += report['weighted avg']['f1-score']
          stop = time.time()
          pAvg = round(precision/runs, 4)
          rAvg = round(recall/runs, 4)
          fAvg = round(f1/runs, 4)
          print("Precision Average :", pAvg)
```

```

print("Recall Average :", rAvg)
print("f1-score Average :", fAvg)
print("Time Elapsed :", round(stop - start, 3), "seconds.")

```

```

Precision Average : 0.6046
Recall Average : 0.4997
f1-score Average : 0.5265
Time Elapsed : 47.019 seconds.

```

From the above Monte-Carlo simulation, we can conclude 2 things - * Both our classifiers have better results than a random guess. * The precision, recall and f1-score values, are close to the $H|T$ probabilities we would observe if we flipped a coin 10,000 times.

1.3 Conclusion

```

In [266]: table = PrettyTable()
          table.field_names = ["Algorithm", "Precision", "Recall", "f1-Score"]
          table.add_row(["Monte Carlo Random Classifier", pAvg, rAvg, fAvg])
          table.add_row(["Logistic Regression Classifier", logitpAvg, logitrAvg, logitfAvg])
          table.add_row(["Random Forest Classifier", rfcAvg, rfcrAvg, rfcfAvg])
          print(table)

```

Algorithm	Precision	Recall	f1-Score
Monte Carlo Random Classifier	0.6046	0.4997	0.5265
Logistic Regression Classifier	0.7262	0.741	0.7312
Random Forest Classifier	0.7036	0.6931	0.6973

In conclusion, it would be apt to say that both our classifiers do a better job than Monte-Carlo Random Classifiers. Both Logistic Regression and Random Forests have precision and recall in the 70 – 80% range. For a small dataset like the one above, these would be satisfactory numbers. That being said, there are opportunities where additional steps can be taken to fine tune the classification model. We can try more creative ways of handling missing data or use down-sampling/up-sampling to better improve our predictions. Data quality and quantity however would play an integral role. More data would help us capture more nuances in the universe. Better data would strengthen our claims.